


Automated Machine Learning (AutoML)

M. Lindauer F. Hutter

University of Freiburg



Lecture 3:

Evaluation and Visualization



Where are we? The big picture

- Introduction
- Background
 - Design spaces in ML
 - Evaluation and visualization
- Hyperparameter optimization (HPO)
 - Bayesian optimization
 - Other black-box techniques
 - Speeding up HPO with multi-fidelity optimization
- Pentecost (Holiday) – no lecture
- Architecture search I + II
- Meta-Learning
- Learning to learn & optimize
- Beyond AutoML: algorithm configuration and control
- Project announcement and closing



After this lecture, you will be able to . . .

- explain the role of outliers in CS/ML
- compare and visualize the performance of different configurations
- compare and visualize the performance of AutoML systems
- explain and correctly apply statistical hypothesis tests



How CS differs from other empirical sciences

- We have a complete and precise **mathematical description** of the object under study
- We have complete and precise **control** of the object under study (and to some degree also the experimental environment)
 - as a result, experiments can be **reproduced perfectly**



How CS differs from other empirical sciences

- We have a complete and precise **mathematical description** of the object under study
- We have complete and precise **control** of the object under study (and to some degree also the experimental environment)
 - as a result, experiments can be **reproduced perfectly**
 - what do we need for experiments to be reproducible? 🙌



How CS differs from other empirical sciences

- We have a complete and precise **mathematical description** of the object under study
- We have complete and precise **control** of the object under study (and to some degree also the experimental environment)
 - as a result, experiments can be **reproduced perfectly**
 - what do we need for experiments to be reproducible?
 - Code & dependencies, inputs, environment (→ VM & cloud)



How CS differs from other empirical sciences

- We have a complete and precise **mathematical description** of the object under study
- We have complete and precise **control** of the object under study (and to some degree also the experimental environment)
 - as a result, experiments can be **reproduced perfectly**
 - what do we need for experiments to be reproducible?
 - Code & dependencies, inputs, environment (→ VM & cloud)
- We often have quite **cheap** experiments
 - price for computers is monotonically decreasing
 - often maximal runtimes of 1h; exception: deep learning (up to a week)
 - compare e.g., experimental physics: 1 week of beam time per year



How CS differs from other empirical sciences

- We have a complete and precise **mathematical description** of the object under study
- We have complete and precise **control** of the object under study (and to some degree also the experimental environment)
 - as a result, experiments can be **reproduced perfectly**
 - what do we need for experiments to be reproducible?
 - Code & dependencies, inputs, environment (→ VM & cloud)
- We often have quite **cheap** experiments
 - price for computers is monotonically decreasing
 - often maximal runtimes of 1h; exception: deep learning (up to a week)
 - compare e.g., experimental physics: 1 week of beam time per year
- We can conduct and analyze experiments fully **automatically**
 - we can gather large amounts of data quickly (e.g., 100 repetitions)
 - but: don't confuse statistical significance and relevance



Outliers are quite different in computer experiments

Is the following statement correct? 🙋

“As usual in other empirical sciences, we (in CS) should take care to **detect and remove outliers** before further analysis.”



Outliers are quite different in computer experiments

Is the following statement correct? 🙋

“As usual in other empirical sciences, we (in CS) should take care to **detect and remove outliers** before further analysis.”

Outliers in CS/ML

- outliers should be investigated closely – why is there an outlier?
- outliers are (hopefully) reproducible – narrow down their reason!

Outliers are quite different in computer experiments

Is the following statement correct? 🙋

“As usual in other empirical sciences, we (in CS) should take care to **detect and remove outliers** before further analysis.”

Outliers in CS/ML

- outliers should be investigated closely – why is there an outlier?
- outliers are (hopefully) reproducible – narrow down their reason!
 - **Algorithm:** When characterizing performance of an optimization algorithm, **outliers with poor values can indicate (deep) local optima**

Outliers are quite different in computer experiments

Is the following statement correct? 🙋

“As usual in other empirical sciences, we (in CS) should take care to **detect and remove outliers** before further analysis.”

Outliers in CS/ML

- outliers should be investigated closely – why is there an outlier?
- outliers are (hopefully) reproducible – narrow down their reason!
 - **Algorithm:** When characterizing performance of an optimization algorithm, **outliers with poor values can indicate (deep) local optima**
 - **Environment:** Outliers can indicate a **problem with the environment** (e.g., file system or network issues)

Outliers are quite different in computer experiments

Is the following statement correct? 🙋

“As usual in other empirical sciences, we (in CS) should take care to **detect and remove outliers** before further analysis.”

Outliers in CS/ML

- outliers should be investigated closely – why is there an outlier?
- outliers are (hopefully) reproducible – narrow down their reason!
 - **Algorithm:** When characterizing performance of an optimization algorithm, **outliers with poor values can indicate (deep) local optima**
 - **Environment:** Outliers can indicate a **problem with the environment** (e.g., file system or network issues)
 - **Datasets:** When characterizing cost across a distribution of datasets, outliers with small values can indicate trivial datasets.

1 Visualization of Configuration Performance


2 Visualization of AutoML Performance

3 Statistical Hypothesis Testing

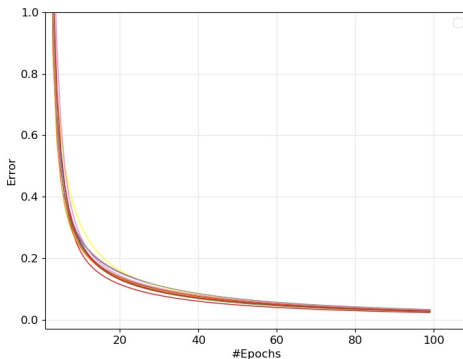
For the following slides, we have used the following setup:

- Model: simple MLP (from sklearn)
with 2 layers with 128 neurons and 64 neurons, resp.
- Dataset: Digits
- Setting 1: learning rate of 0.001
- Setting 2: learning rate of 0.01

A Single Learning Curve (Setting 1)




10 Learning Curves (Setting 1)



- Deep learning (and most other ML algorithms) are non-deterministic
- Measure performance more than once and estimate noise-level




Aggregated Learning Curves (Setting 1)



- Plot mean and stdev (shaded area) across n (here 100) random seeds
- Only use mean and stdev if noise is somehow normal distributed
- alternatives are the mean+standard error or median+25/75-percentiles


Comparing Learning Curves



- If uncertainties (shaded area) overlap, results might not be statistically significant but due to noise

Scatter Plot: Train vs. Test


Setting 1



- perfect would be if train and test score are correlated
- here at least not anti-correlated


Scatter Plot: Train vs. Test

Setting 1



- perfect would be if train and test score are correlated
- here at least not anti-correlated


Setting 2



- already perfect training score
- generalization nevertheless noisy




eCDF: Distribution of Performance (Setting 1)



How to compute an empirical CDF:


- 1 X: sorted error scores
- 2 Y: $[1 \dots \text{\#points}] / \text{\#points}$
- 3 `step_function(X,Y)`

eCDF: Distribution of Performance




Which curve corresponds to which setting? 🙌

Box Plot: Comparing Two Configurations



- alternative: violin plots

Box Plot: Comparing Two Configurations



- alternative: violin plots
 - if you have paired populations (e.g., configurations evaluated on different datasets), you should generate boxplots with $\mathcal{L}/\mathcal{L}_{\text{baseline}}$
- ~> insight on how many datasets one of the two performed better



- 1 Visualization of Configuration Performance
- 2 Visualization of AutoML Performance
- 3 Statistical Hypothesis Testing

- Don't only measure the performance of AutoML systems for a fixed budget because:
 - ① A priori, it is unknown how long a user will run an AutoML system
 - ② Some systems perform well for small budgets and some others need more time before performing well
 - BO-based systems often perform as good as random search in the beginning, but perform very well later on



AutoML Systems over Time

- Don't only measure the performance of AutoML systems for a fixed budget because:
 - ① A priori, it is unknown how long a user will run an AutoML system
 - ② Some systems perform well for small budgets and some others need more time before performing well
 - BO-based systems often perform as good as random search in the beginning, but perform very well later on
- ~> AutoML systems should have a good anytime performance




AutoML Systems over Time

- Don't only measure the performance of AutoML systems for a fixed budget because:
 - ① A priori, it is unknown how long a user will run an AutoML system
 - ② Some systems perform well for small budgets and some others need more time before performing well
 - BO-based systems often perform as good as random search in the beginning, but perform very well later on
- ↪ AutoML systems should have a good anytime performance
- Recommendation: plot performance (e.g., error) over time
 - Similar to learning curves of DNNs



Aggregated AutoML Systems over Time



- Different HPO systems minimizing the error of an MLP on MNIST
- How to do it:
 - Run each systems several times and always log its current incumbent
 - Plot mean and stdev for each system after each incumbent update
 - Use step functions, because linear interpolation would be too optimistic



Aggregated AutoML Systems over Time and Datasets

- An AutoML system should not only perform well on a single dataset but on many datasets
- To summarize and compare the performance across a set of datasets, we can plot the performance over time and across datasets



Aggregated AutoML Systems over Time and Datasets

- An AutoML system should not only perform well on a single dataset but on many datasets
- To summarize and compare the performance across a set of datasets, we can plot the performance over time and across datasets
- Ways to do it:
 - 1 Average loss metric (e.g., error) across datasets
 - Problem: different scales of errors on different datasets



Aggregated AutoML Systems over Time and Datasets

- An AutoML system should not only perform well on a single dataset but on many datasets
- To summarize and compare the performance across a set of datasets, we can plot the performance over time and across datasets
- Ways to do it:
 - 1 Average loss metric (e.g., error) across datasets
 - Problem: different scales of errors on different datasets
 - 2 Normalize loss metric across datasets, e.g., best seen loss relates to a zero cost
 - Problem: Minor improvements ($\leq 0.1\%$) might look substantial




Aggregated AutoML Systems over Time and Datasets

- An AutoML system should not only perform well on a single dataset but on many datasets
 - To summarize and compare the performance across a set of datasets, we can plot the performance over time and across datasets
 - Ways to do it:
 - 1 Average loss metric (e.g., error) across datasets
 - Problem: different scales of errors on different datasets
 - 2 Normalize loss metric across datasets, e.g., best seen loss relates to a zero cost
 - Problem: Minor improvements ($\leq 0.1\%$) might look substantial
- ~> There is no way around that you have some kind of information loss

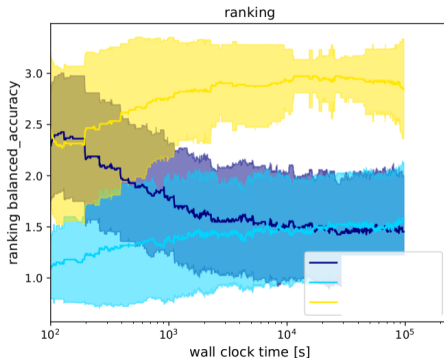


Ranked AutoML Systems over Time and Datasets



- Remark: x-axis should be on log-scale

Ranked AutoML Systems over Time and Datasets



- Remark: x-axis should be on log-scale
- Ranks avoid the problem of different scales
- Problem: we don't know whether a better ranking relates to a substantial improvement of the actual cost metric

Uncertainties in Ranking Plots

Given:

- k systems you want compare
- n runs of each system
- d datasets



Uncertainties in Ranking Plots

Given:

- k systems you want compare
- n runs of each system
- d datasets

Across Datasets ($d > 1$)

- 1 for each system and dataset, average performance across the n runs
- 2 compute ranking on each datasets
- 3 compute mean and stdev across rankings



Uncertainties in Ranking Plots

Given:

- k systems you want compare
- n runs of each system
- d datasets


Across Datasets ($d > 1$)

- 1 for each system and dataset, average performance across the n runs
- 2 compute ranking on each datasets
- 3 compute mean and stdev across rankings

On a single dataset ($d = 1$)

- 1 for each system, sample a single run out of the n runs
- 2 compute ranking on these samples
- 3 repeat 1. and 2. at least 10 000 times to obtain many rankings
- 4 compute mean and stdev across rankings


Scatter Plot: Meta Features



- Compare two systems (or 1 vs rest)
- each dot is a dataset
- color encoding better or worse
- size indicates performance difference
- meta-features (e.g., #features, #samples) on axes
 - using PCA on a larger set of meta-features leads to algorithm footprints [Smith-Miles et al. 2012]



Heatmaps: Comparing several Configurations



- compare many configurations on many datasets
- can reveal:
 - homogeneity if smooth transition from hard to easy
 - heterogeneity if stripes exist (here the case)
- Remark: datasets and configurations should be sorted according to their average cost

- 1 Visualization of Configuration Performance
- 2 Visualization of AutoML Performance
- 3 Statistical Hypothesis Testing

Background: statistical hypothesis tests

- When we have a lot of data, we need to summarize it
 - But we already saw that summarization hides a lot of data
 - Ideally, we want to draw high-level conclusions (e.g., “A outperforms B on datasets of type X”)



Background: statistical hypothesis tests

- When we have a lot of data, we need to summarize it
 - But we already saw that summarization hides a lot of data
 - Ideally, we want to draw high-level conclusions (e.g., “A outperforms B on datasets of type X”)
 - Problem: we only have a finite number of observations
 - Can we attribute observed performance differences to chance?
 - Are we reasonably sure that a claim we make is reproducible?
- ~> Statistical tests can help



- 1 Define initial research hypothesis

- 1 Define initial research hypothesis
- 2 Derive null H_0 and alternative H_1 hypothesis
 - Alternative hypothesis should be your research hypothesis

First example: Courtroom Trial

- A prosecutor tries to prove the guilt of the defendant
- H_0 : 🙌
- H_1 : 🙌

First example: Courtroom Trial

- A prosecutor tries to prove the guilt of the defendant
- H_0 : ✎ The defendant is not guilty
 - Accepted for the moment
(“not guilty as long as their guilt is not proven”)
- H_1 : ✎ The defendant is guilty
 - prosecutor hopes to support that



First example: Courtroom Trial

- A prosecutor tries to prove the guilt of the defendant
- H_0 : ✎ The defendant is not guilty
 - Accepted for the moment
(“not guilty as long as their guilt is not proven”)
- H_1 : ✎ The defendant is guilty
 - prosecutor hopes to support that

	Truly not guilty	Truly guilty
Found not guilty	Acquittal	Type II Error
Found guilty	Type I Error	Conviction

⇒ We want to minimize Type I error!



Statistical hypothesis testing (cont'd)

- 1 Define initial research hypothesis
- 2 Derive null H_0 and alternative H_1 hypothesis
 - Alternative hypothesis should be your research hypothesis
- 3 Consider statistical assumptions
 - E.g., is your data Gaussian distributed?



Statistical hypothesis testing (cont'd)

- ① Define initial research hypothesis
- ② Derive null H_0 and alternative H_1 hypothesis
 - Alternative hypothesis should be your research hypothesis
- ③ Consider statistical assumptions
 - E.g., is your data Gaussian distributed?
- ④ Decide test and test statistic T
 - The correct test depends on your statistical assumptions.
 - Typically: if you use more assumptions, the test is more powerful (i.e., less Type-I error)



Statistical hypothesis testing (cont'd)

- ① Define initial research hypothesis
- ② Derive null H_0 and alternative H_1 hypothesis
 - Alternative hypothesis should be your research hypothesis
- ③ Consider statistical assumptions
 - E.g., is your data Gaussian distributed?
- ④ Decide test and test statistic T
 - The correct test depends on your statistical assumptions.
 - Typically: if you use more assumptions, the test is more powerful (i.e., less Type-I error)
- ⑤ Decide significance level α
(i.e., acceptable Type-I error to reject null hypothesis)



Statistical hypothesis testing (cont'd)

- ① Define initial research hypothesis
- ② Derive null H_0 and alternative H_1 hypothesis
 - Alternative hypothesis should be your research hypothesis
- ③ Consider statistical assumptions
 - E.g., is your data Gaussian distributed?
- ④ Decide test and test statistic T
 - The correct test depends on your statistical assumptions.
 - Typically: if you use more assumptions, the test is more powerful (i.e., less Type-I error)
- ⑤ Decide significance level α
(i.e., acceptable Type-I error to reject null hypothesis)
- ⑥ Compute observed t_{obs} of test statistic T
- ⑦ Calculate p -value given t_{obs}
 - i.e., probability under the null hypothesis of sampling a test statistic as extreme as observed (probability of Type-I error)



Statistical hypothesis testing (cont'd)

- 1 Define initial research hypothesis
- 2 Derive null H_0 and alternative H_1 hypothesis
 - Alternative hypothesis should be your research hypothesis
- 3 Consider statistical assumptions
 - E.g., is your data Gaussian distributed?
- 4 Decide test and test statistic T
 - The correct test depends on your statistical assumptions.
 - Typically: if you use more assumptions, the test is more powerful (i.e., less Type-I error)
- 5 Decide significance level α
(i.e., acceptable Type-I error to reject null hypothesis)
- 6 Compute observed t_{obs} of test statistic T
- 7 Calculate p -value given t_{obs}
 - i.e., probability under the null hypothesis of sampling a test statistic as extreme as observed (probability of Type-I error)
- 8 If $p < \alpha$, reject null hypothesis in favor of alternative hypothesis
 - If $p > \alpha$, it doesn't tell you anything about the null hypothesis!



Second example for a statistical test

- Claim: “the students in this class are more intelligent than average”



Second example for a statistical test

- Claim: “the students in this class are more intelligent than average”
- Null Hypothesis H_0 : $\mu = 100$ (μ is the population mean of this class)
- Alternative Hypothesis H_1 : $\mu > 100$ (one-sided hypothesis)



Second example for a statistical test

- Claim: “the students in this class are more intelligent than average”
- Null Hypothesis H_0 : $\mu = 100$ (μ is the population mean of this class)
- Alternative Hypothesis H_1 : $\mu > 100$ (one-sided hypothesis)
- IQ values are known to be normally distributed with $X \sim \mathcal{N}(100, 15)$
 - \rightarrow statistical assumption



Second example for a statistical test

- Claim: “the students in this class are more intelligent than average”
- Null Hypothesis H_0 : $\mu = 100$ (μ is the population mean of this class)
- Alternative Hypothesis H_1 : $\mu > 100$ (one-sided hypothesis)
- IQ values are known to be normally distributed with $X \sim \mathcal{N}(100, 15)$
 - \rightarrow statistical assumption
- Let's say we observed IQ values x_i of 9 students in the class:
 - $\{x_1, \dots, x_9\} = \{116, 128, 125, 119, 89, 99, 105, 116, 118\}$.
 - The sample mean is $\bar{x} = 112.8$
 - Does this data support the claim?



Example continued

- Distribution of the test statistic


- Under H_0 , we know that each $x_i \sim \mathcal{N}(100, 15)$
- The test statistic that we measure is the sample mean $\bar{x} = \frac{1}{9} \sum_{i=1}^9 x_i$




Example continued

- Distribution of the test statistic

- Under H_0 , we know that each $x_i \sim \mathcal{N}(100, 15)$
- The test statistic that we measure is the sample mean $\bar{x} = \frac{1}{9} \sum_{i=1}^9 x_i$
- Under H_0 , the distribution of \bar{x} is $\mathcal{N}(100, 15/\sqrt{9})$
 - Our observation $\bar{x} = 112.8$ is quite extreme under that distribution




General principle




- Compare the test statistic (here: \bar{x}) to its sampling distribution under H_0

General principle




- Compare the test statistic (here: \bar{x}) to its sampling distribution under H_0
- **P-value**: probability p of observing values at least as extreme as \bar{x}

General principle



- Compare the test statistic (here: \bar{x}) to its sampling distribution under H_0
- **P-value**: probability p of observing values **at least as extreme as \bar{x}**
- Compare p to pre-defined confidence level α (usually $\alpha = 0.05$);
if $p < \alpha$, **reject H_0**

General principle



- Compare the test statistic (here: \bar{x}) to its sampling distribution under H_0
- **P-value**: probability p of observing values **at least as extreme as \bar{x}**
- Compare p to pre-defined confidence level α (usually $\alpha = 0.05$);
if $p < \alpha$, **reject H_0**
- With $\alpha = 0.01$, would we reject H_0 in this case? 🙌

Summary of example

- We just used a so-called *Z-test*
- $H_0: \mu = \mu_0, H_1: \mu > \mu_0$
- Assumptions: $X \sim \mathcal{N}(\mu, \sigma^2)$, with known μ and σ^2



Summary of example

- We just used a so-called **Z-test**
- $H_0: \mu = \mu_0, H_1: \mu > \mu_0$
- Assumptions: $X \sim \mathcal{N}(\mu, \sigma^2)$, with known μ and σ^2
- **Test statistic:** sample mean \bar{x} ; evaluate under $\mathcal{N}(\mu = \mu_0, s = \sigma^2 / \sqrt{n})$



Summary of example

- We just used a so-called **Z-test**
- $H_0: \mu = \mu_0, H_1: \mu > \mu_0$
- Assumptions: $X \sim \mathcal{N}(\mu, \sigma^2)$, with known μ and σ^2
- **Test statistic**: sample mean \bar{x} ; evaluate under $\mathcal{N}(\mu = \mu_0, s = \sigma^2/\sqrt{n})$
- Equivalent: compute the **Z-statistic**: $Z = (\bar{x} - \mu_0)/s$ and evaluate cumulative distribution $\Phi(Z)$ of Z under $\mathcal{N}(0, 1)$




Summary of example

- We just used a so-called **Z-test**
- $H_0: \mu = \mu_0, H_1: \mu > \mu_0$
- Assumptions: $X \sim \mathcal{N}(\mu, \sigma^2)$, with known μ and σ^2
- **Test statistic**: sample mean \bar{x} ; evaluate under $\mathcal{N}(\mu = \mu_0, s = \sigma^2/\sqrt{n})$
- Equivalent: compute the **Z-statistic**: $Z = (\bar{x} - \mu_0)/s$ and evaluate cumulative distribution $\Phi(Z)$ of Z under $\mathcal{N}(0, 1)$
 - There are standard tables to look up $\Phi(Z)$ for different values of Z
 - Nowadays, there are standard libraries to compute $\Phi(Z)$




Two-sided tests




- Similar to one-sided tests, but testing for extreme values in both tails
- Example Z-test: two-sided alternative hypothesis $H_1: \mu \neq \mu_0$

Two-sided tests



- Similar to one-sided tests, but testing for extreme values in both tails
- Example Z-test: two-sided alternative hypothesis $H_1: \mu \neq \mu_0$
- Compute $Z = (\bar{x} - \mu_0)/s$ as before
- Compute p -value as $p = 2\Phi(Z)$, to account for both tails

Two-sided tests



- Similar to one-sided tests, but testing for extreme values in both tails
- Example Z-test: two-sided alternative hypothesis $H_1: \mu \neq \mu_0$
- Compute $Z = (\bar{x} - \mu_0)/s$ as before
- Compute p -value as $p = 2\Phi(Z)$, to account for both tails
- With $\alpha = 0.01$, would we reject H_0 in this case? 🙌

General points about statistical hypothesis tests

- What if $p > \alpha$?
 - Failure to reject H_0
 - This does not mean that we accept H_0



General points about statistical hypothesis tests

- What if $p > \alpha$?
 - Failure to reject H_0
 - This does not mean that we accept H_0
- Beware (i): most tests make some assumptions
 - E.g., Z -test and popular t -test assume normality
 - Our data is often far from normally-distributed
 - ~> E.g., exponential runtime distributions of local search optimizers
 - ~> E.g., distribution of fitting a neural network with different random seeds is not well studied



General points about statistical hypothesis tests

- What if $p > \alpha$?
 - Failure to reject H_0
 - This does not mean that we accept H_0
- Beware (i): most tests make some assumptions
 - E.g., Z -test and popular t -test assume **normality**
 - Our data is often far from normally-distributed
 - ~> E.g., exponential runtime distributions of local search optimizers
 - ~> E.g., distribution of fitting a neural network with different random seeds is not well studied
- Beware (ii): if you use cross-validation, observations are not independent (you cannot apply statistical tests that assume independence)



The Wilcoxon rank-sum test for non-normal data

- Compare the distributions of random variables X and Y
 - based on samples x_1, \dots, x_n and y_1, \dots, y_m
- Assumptions
 - All of x_1, \dots, x_n and y_1, \dots, y_m are independent from each other
 - Responses are ordinal (we can compare them)



The Wilcoxon rank-sum test for non-normal data

- Compare the distributions of random variables X and Y
 - based on samples x_1, \dots, x_n and y_1, \dots, y_m
- Assumptions
 - All of x_1, \dots, x_n and y_1, \dots, y_m are independent from each other
 - Responses are ordinal (we can compare them)
- $H_0: P(X > Y) = P(Y > X)$



The Wilcoxon rank-sum test for non-normal data

- Compare the distributions of random variables X and Y
 - based on samples x_1, \dots, x_n and y_1, \dots, y_m
- Assumptions
 - All of x_1, \dots, x_n and y_1, \dots, y_m are independent from each other
 - Responses are ordinal (we can compare them)
- $H_0: P(X > Y) = P(Y > X)$
- $H_1: P(X > Y) > P(Y > X)$ (one-sided)



The Wilcoxon rank-sum test for non-normal data

- Compare the distributions of random variables X and Y
 - based on samples x_1, \dots, x_n and y_1, \dots, y_m
- Assumptions
 - All of x_1, \dots, x_n and y_1, \dots, y_m are independent from each other
 - Responses are ordinal (we can compare them)
- $H_0: P(X > Y) = P(Y > X)$
- $H_1: P(X > Y) > P(Y > X)$ (one-sided)
- Test statistic
 - Order all elements x_i and y_j and give them ranks (1 for smallest)
 - Compute the sum of ranks of x_i and of y_j



The Wilcoxon rank-sum test for non-normal data

- Compare the distributions of random variables X and Y
 - based on samples x_1, \dots, x_n and y_1, \dots, y_m
- Assumptions
 - All of x_1, \dots, x_n and y_1, \dots, y_m are independent from each other
 - Responses are ordinal (we can compare them)
- $H_0: P(X > Y) = P(Y > X)$
- $H_1: P(X > Y) > P(Y > X)$ (one-sided)
- Test statistic
 - Order all elements x_i and y_j and give them ranks (1 for smallest)
 - Compute the sum of ranks of x_i and of y_j
- Under H_0 , the distribution of that test statistic is known
 \rightsquigarrow can evaluate how extreme the observed test statistic is



The permutation test: another test for non-normal data

- Framework for testing several types of claims
- E.g., H_0 : X and Y have equal means
- Test statistic: $t = \frac{1}{n} \sum_{i=1}^n x_i - \frac{1}{m} \sum_{j=1}^m y_j$



The permutation test: another test for non-normal data

- Framework for testing several types of claims
- E.g., H_0 : X and Y have equal means
- Test statistic: $t = \frac{1}{n} \sum_{i=1}^n x_i - \frac{1}{m} \sum_{j=1}^m y_j$
- The sampling distribution to compare t against:
 - Put $\{x_1, \dots, x_n\}$ and $\{y_1, \dots, y_m\}$ into a single pool
 - $S = []$; repeat, e.g., 10 000 times
 - draw a random permutation & permute pool with it
 - add test statistic over permuted pool to S



The permutation test: another test for non-normal data

- Framework for testing several types of claims
- E.g., H_0 : X and Y have equal means
- Test statistic: $t = \frac{1}{n} \sum_{i=1}^n x_i - \frac{1}{m} \sum_{j=1}^m y_j$
- The sampling distribution to compare t against:
 - Put $\{x_1, \dots, x_n\}$ and $\{y_1, \dots, y_m\}$ into a single pool
 - $S = []$; repeat, e.g., 10 000 times
 - draw a random permutation & permute pool with it
 - add test statistic over permuted pool to S
- p -value: percentile of s in S :
fraction of samples s in S with $s < t$



Paired vs. unpaired tests

- if you have two unsorted populations (e.g., repeated measurements with different random seeds), you use an unpaired test (as discussed above)
- if you can attribute a measurement to concrete objects (e.g., measurements on different datasets), you use a paired test
 - paired test are more powerful
(i.e., higher confidence for the same amount of data)
- Examples for paired permutation tests
 - Wilcoxon signed rank test
 - paired permutation test
 - permutation of measurement pairs (e.g., same dataset)



Paired Permutation Test

Given:

- **ordered** observations $[x_1, \dots, x_n]$ in X (resp. in Y) where each observation is attributed to concrete objects



Paired Permutation Test

Given:

- **ordered** observations $[x_1, \dots, x_n]$ in X (resp. in Y) where each observation is attributed to concrete objects
- $S = []$; repeat, e.g., 10 000 times
 - 1 $X' = []$ and $Y' = []$
 - 2 for each **pair of observations** $\langle x_i, y_i \rangle$ sample whether (i) you put x_i into X' and y_i into Y' or (ii) x_i into Y' and y_i into X' .
 - 3 add test statistic based on X' and Y' to S



Multiple Testing Correction

- To compare many systems, we apply statistical tests several time (once for each pair of systems)
- Each test induces some error (bounded by α)

⇒ the error probability for k tests is bounded by $1 - (1 - \alpha)^k$

k	1	2	3	4	5	6	7
$1 - (1 - 0.05)^k$	0.05	0.10	0.14	0.19	0.23	0.27	0.30



Multiple Testing Correction

- To compare many systems, we apply statistical tests several time (once for each pair of systems)
- Each test induces some error (bounded by α)

~> the error probability for k tests is bounded by $1 - (1 - \alpha)^k$

k	1	2	3	4	5	6	7
$1 - (1 - 0.05)^k$	0.05	0.10	0.14	0.19	0.23	0.27	0.30

- Better would be if the error probability would not increase with k

~> multiple testing correction



Multiple Testing Correction

- To compare many systems, we apply statistical tests several time (once for each pair of systems)
- Each test induces some error (bounded by α)

~> the error probability for k tests is bounded by $1 - (1 - \alpha)^k$

k	1	2	3	4	5	6	7
$1 - (1 - 0.05)^k$	0.05	0.10	0.14	0.19	0.23	0.27	0.30

- Better would be if the error probability would not increase with k

~> multiple testing correction

- Bonferroni testing correction: $\alpha_{\text{local}} = \alpha_{\text{global}}/k$



Multiple Testing Correction

- To compare many systems, we apply statistical tests several time (once for each pair of systems)
- Each test induces some error (bounded by α)

~> the error probability for k tests is bounded by $1 - (1 - \alpha)^k$

k	1	2	3	4	5	6	7
$1 - (1 - 0.05)^k$	0.05	0.10	0.14	0.19	0.23	0.27	0.30

- Better would be if the error probability would not increase with k

~> multiple testing correction

- Bonferroni testing correction: $\alpha_{\text{local}} = \alpha_{\text{global}}/k$
 - very conservative approach
 - there exist other, less conservative approaches



Checklist for good scientific practices

Incomplete list of good scientific practices (specifically for students):

- ① keep track of your code and design decisions (on all levels)
- ② measure performance of randomized algorithms multiple times and show uncertainty of results
- ③ apply suitable statistical tests to check for significance
- ④ choose a metric that is relevant for the application
- ⑤ always add legends, axis labels and so on in plots
- ⑥ be aware of other research results
- ⑦ avoid peeking at your test data
 - no cherry-picking!



Now, you should be able to ...

- explain the role of outliers in CS/ML
- compare and visualize the performance of different configurations
- compare and visualize the performance of AutoML systems
- explain and correctly apply statistical hypothesis tests

Literature [These are links]

- [P. Langley 1988. Machine Learning as an Experimental Science]
- [C. Drummod 2006. Machine Learning as an Experimental Science (Revisited)]
- [J. Demsar 2006. Statistical Comparisons of Classifiers over Multiple Data Sets]
- [Wilson et al. 2014. Best Practices for Scientific Computing]
- [Wilson et al. 2017. Good enough practices in scientific computing]

