

Multi-criteria Optimization

Bayesian Optimization

Bernd Bischl Frank Hutter Lars Kotthoff
Marius Lindauer Joaquin Vanschoren

Recap: Bayesian Optimization I

Advantages of BO

- Sample efficient
- Can handle noise
- Native incorporation of priors
- Does not require gradients
- Theoretical guarantees

We will now extend BO to multiple cost functions.

Recap: Bayesian Optimization II

Bayesian optimization loop

Require: Search space Λ , cost function c , acquisition function u , predictive model \hat{c} , maximal number of function evaluations T

Result : Best configuration $\hat{\lambda}$ (according to \mathcal{D} or \hat{c})

- 1 Initialize data $\mathcal{D}^{(0)}$ with initial observations
 - 2 **for** $t = 1$ **to** T **do**
 - 3 Fit predictive model $\hat{c}^{(t)}$ on $\mathcal{D}^{(t-1)}$
 - 4 Select next query point: $\lambda^{(t)} \in \arg \max_{\lambda \in \Lambda} u(\lambda; \mathcal{D}^{(t-1)}, \hat{c}^{(t)})$
 - 5 Query $c(\lambda^{(t)})$
 - 6 Update data: $\mathcal{D}^{(t)} \leftarrow \mathcal{D}^{(t-1)} \cup \{(\lambda^{(t)}, c(\lambda^{(t)}))\}$
-

Multi-Criteria Bayesian Optimization

Goal: Extend Bayesian optimization to multiple cost functions

$$\min_{\lambda \in \Lambda} c(\lambda) \Leftrightarrow \min_{\lambda \in \Lambda} (c_1(\lambda), c_2(\lambda), \dots, c_m(\lambda)) .$$

There are two basic approaches:

- 1 Simplify the problem by scalarizing the cost functions, or
- 2 define acquisition functions for multiple cost functions.

Scalarization

Idea: Aggregate all cost functions

$$\min_{\lambda \in \Lambda} \sum_{i=1}^m w_i c_i(\lambda) \quad \text{with} \quad w_i \geq 0$$

- **Obvious problem:** How to choose w_1, \dots, w_m ?
 - ▶ Expert knowledge?
 - ▶ Systematic variation?
 - ▶ Random variation?
- If expert knowledge is not available a-priori, we need to ensure that different trade-offs between cost functions are explored.
- Simplifies multi-criteria optimization problem to single-objective
 - Bayesian optimization can be used without adaption of the general algorithm.

Scalarize the cost functions using the augmented Tchebycheff norm / achievement function

$$c = \max_{i=1,\dots,m} (w_i c_i(\boldsymbol{\lambda})) + \rho \sum_{i=1}^m w_i c_i(\boldsymbol{\lambda}),$$

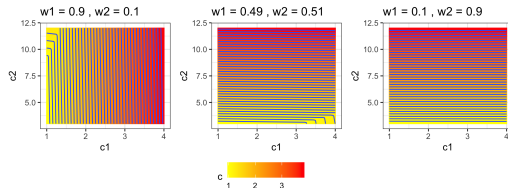
- The weights $w \in W$ are drawn from

$$W = \left\{ w = (w_1, \dots, w_m) \mid \sum_{i=1}^m w_i = 1, w_i = \frac{l}{s} \wedge, l \in 0, \dots, s \right\},$$

with $|W| = \binom{s+m-1}{k-1} 1$.

- New weights are drawn in every BO iteration.
- ρ is a small parameter suggested to be set to 0.05.
- s selects the number of different weights to draw from.

Why the Tchebycheff norm?



$$c = \max_{i=1, \dots, m} (w_i c_i(\lambda)) + \rho \sum_{i=1}^m w_i c_i(\lambda),$$

- The norm consists of two components:
 - ▶ $\max_{i=1, \dots, m} (w_i c_i(\lambda))$ takes only the maximum weighted cost into account.
 - ▶ $\sum_{i=1}^m w_i c_i(\lambda)$ is the weighted sum of all cost functions.
- ρ describes the trade-off between these components.
- By the randomized weights in each iteration and the usually small value of $\rho = 0.05$, this allows exploration of extreme points of single cost functions.
- One can prove: **Every solution of the scalarized problem is pareto-optimal!**

ParEGO Algorithm

ParEGO loop

Require: Search space Λ , cost function c , acquisition function u , predictive model \hat{c} , maximal number of function evaluations T , ρ , l , s

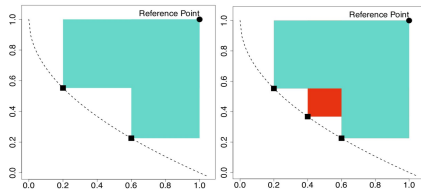
Result : Best configuration $\hat{\lambda}$ (according to \mathcal{D} or \hat{c})

- 1 Initialize data $\mathcal{D}^{(0)}$ with initial observations
 - 2 **for** $t = 1$ **to** T **do**
 - 3 Sample w from $\{w = (w_1, \dots, w_m) \mid \sum_{i=1}^m w_i = 1, w_i = \frac{l}{s} \wedge, l \in 0, \dots, s\}$;
 - 4 Compute scalarization $c^{(t)} = \max_{i=1, \dots, m} (w_i c_i(\lambda)) + \rho \sum_{i=1}^m w_i c_i(\lambda)$;
 - 5 Fit predictive model $\hat{c}^{(t)}$ on $\mathcal{D}^{(t-1)}$
 - 6 Select next query point: $\lambda^{(t)} \in \arg \max_{\lambda \in \Lambda} u(\lambda; \mathcal{D}^{(t-1)}, \hat{c}^{(t)})$
 - 7 Query $c(\lambda^{(t)})$
 - 8 Update data: $\mathcal{D}^{(t)} \leftarrow \mathcal{D}^{(t-1)} \cup \{\langle \lambda^{(t)}, c(\lambda^{(t)}) \rangle\}$
-

Hypervolume based Acquisition Functions

Idea: Define acquisition function that directly models contribution to dominated HV.

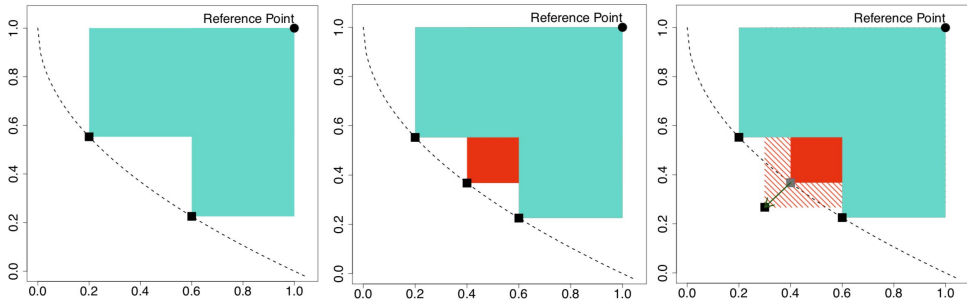
$$\max(0, S(\mathcal{P} \cup \boldsymbol{\lambda}, R) - S(\mathcal{P}, R))$$



- Fit m single-objective surrogate models $\hat{c}_1, \dots, \hat{c}_m$
- Acquisition function takes all surrogate models into account.
- Single-criteria optimization of acquisition function.

S-Metric Selection-based EGO I

Using the Lower Confidence bound $u_{\text{LCB},1}(\lambda), \dots, u_{\text{LCB},m}(\lambda)$, an optimistic estimate of hypervolume contribution can be calculated.



S-Metric Selection-based EGO II

Problem: Based on the way the hypervolume contribution is measured large plateaus of zero improvement are present.

- These make optimization much harder.
- An adaptive penalty is added to regions in which the lower confidence bound is dominated.

This method is referred to as SMS-EGO [Ponweiser et al. 2008].

Further Hypervolume based Acquisition Functions

Expected Hypervolume Improvement (EHI) [Yang et al. 2019]

$$u_{EI, \mathcal{H}}(\boldsymbol{\lambda}) = \int_{-\infty}^{\infty} p(c \mid \boldsymbol{\lambda}) \times \mathcal{H}(\boldsymbol{\lambda}) \, dc,$$

with $\mathcal{H}(\boldsymbol{\lambda}) = S(\mathcal{P} \cup \boldsymbol{\lambda}, R) - S(\mathcal{P}, R)$.

- Direct extension of u_{EI} to the hypervolume.
- $p(c \mid \boldsymbol{\lambda})$ is the joint density of the surrogate model predictions at $\boldsymbol{\lambda}$.
- As the surrogates are GPs and modeled independently of each other, this is just an integral over m univariate normal distributions.
- Efficient computations for $m \leq 3$ exist, beyond that expensive simulation-based computation is required.

Further hypervolume based acquisition functions:

- **Stepwise Uncertainty Reduction** (SUR) based on the probability of improvement.
- **Expected Maximin Improvement** (EMI) based on the ϵ -indicator.

Hypervolume based BO Algorithm

Hypervolume based Bayesian optimization loop

Require: Search space Λ , cost function c , acquisition function u , predictive model \hat{c} , maximal number of function evaluations T

Result : Best configuration $\hat{\lambda}$ (according to \mathcal{D} or \hat{c})

- 1 Initialize data $\mathcal{D}^{(0)}$ with initial observations
 - 2 **for** $t = 1$ **to** T **do**
 - 3 Fit predictive models $\hat{c}_1^{(t)}, \dots, \hat{c}_m^{(t)}$ on $\mathcal{D}^{(t-1)}$
 - 4 Select next query point: $\lambda^{(t)} \in \arg \max_{\lambda \in \Lambda} u(\lambda; \mathcal{D}^{(t-1)}, \hat{c}_1^{(t)}, \dots, \hat{c}_m^{(t)})$
 - 5 Query $c(\lambda^{(t)})$
 - 6 Update data: $\mathcal{D}^{(t)} \leftarrow \mathcal{D}^{(t-1)} \cup \{\langle \lambda^{(t)}, c(\lambda^{(t)}) \rangle\}$
-

Multi-criteria Optimization

Introduction

Bernd Bischl Frank Hutter Lars Kotthoff
Marius Lindauer Joaquin Vanschoren

Introductory example I

Often we want to solve optimization problems concerning several goals.

General applications:

- Medicine: maximum effect, but minimum side effect of a drug.
- Finances: maximum return, but minimum risk of an equity portfolio.
- Production planning: maximum revenue, but minimum costs.
- Booking a hotel: maximum rating, but minimum costs.

In machine learning:

- Sparse models: maximum predictive performance, but minimal number of features.
- Fast models: maximum predictive performance, but short prediction time.
- ...

Introductory example II

Example:

Choose the best hotel to stay at by maximizing ratings subject to a maximum price per night.

Problems:

- The result depends on how we select the maximum price and usually returns different solutions for different maximum price values.
- We could also choose a minimum rating and optimize the price per night.
- The more objectives we optimize, the more difficult such a definition becomes.

Goal:

Find a more general approach to solve multi-criteria problems.

Introductory example III




Maritim Hotel München

4,0 ★★★★★ (899)

📶 Kostenloses WLAN

76 €



H+ Hotel München

4,2 ★★★★★ (660)

📶 Kostenloses WLAN

84 €




Marriott Hotel München

4,3 ★★★★★ (1.030)

28 % Rabatt

107 €
77 €



Hotel Vier Jahreszeiten
Kempinski Munich

4,6 ★★★★★ (1.025)

📶 Kostenloses WLAN

278 €

When booking a hotel: find the hotel with

- minimum price per night (**costs**) and
- maximum user rating (**performance**).


Since our standard is to minimize objectives, we minimize negative ratings.

Introductory example IV

The objectives often conflict with each other:

- Lower price \rightarrow usually lower hotel rating.
- Better rating \rightarrow usually higher price.



Example: (negative) average rating by hotel guests (1 - 5) vs. average price per night (excerpt).



Introductory example V

Often, objectives are not directly comparable as they are measured on different scales:

- Left: A hotel with rating 4 for 89 Euro ($c^{(1)} = (89, -4.0)$) would be preferred to a hotel for 108 Euro with the same rating ($c^{(2)} = (108, -4.0)$).
- Right: How to decide if $c^{(1)} = (89, -4.0)$ or $c^{(1)} = (95, -4.5)$ is preferred?
- How much is one *rating point* worth?



Definition: multi-criteria optimization problem

A **multi-criteria optimization problem** is defined by

$$\min_{\boldsymbol{\lambda} \in \boldsymbol{\Lambda}} c(\boldsymbol{\lambda}) \Leftrightarrow \min_{\boldsymbol{\lambda} \in \boldsymbol{\Lambda}} (c_1(\boldsymbol{\lambda}), c_2(\boldsymbol{\lambda}), \dots, c_m(\boldsymbol{\lambda})),$$

with $\boldsymbol{\Lambda} \subset \mathbb{R}^n$ and multi-criteria objective function $c : \boldsymbol{\Lambda} \rightarrow \mathbb{R}^m$, $m \geq 2$.

- **Goal:** minimize multiple target functions simultaneously.
- $(c_1(\boldsymbol{\lambda}), \dots, c_m(\boldsymbol{\lambda}))^\top$ maps each candidate $\boldsymbol{\lambda}$ into the objective space \mathbb{R}^m .
- Often no clear best solution, as objective are usually conflicting and we cannot totally order in \mathbb{R}^m .
- W.l.o.g. we always minimize.
- Alternative names: multi-criteria optimization, multi-objective optimization, Pareto optimization.

Pareto sets and Pareto optimality

Definition:

Given a multi-criteria optimization problem

$$\min_{\lambda \in \Lambda} (c_1(\lambda), \dots, c_m(\lambda)), \quad c_i : \Lambda \rightarrow \mathbb{R}.$$


- A candidate $\lambda^{(1)}$ (**Pareto-**) **dominates** $\lambda^{(2)}$, if $c(\lambda^{(1)}) \prec c(\lambda^{(2)})$, i.e.
 - ① $c_i(\lambda^{(1)}) \leq c_i(\lambda^{(2)})$ for all $i \in \{1, 2, \dots, m\}$ and
 - ② $c_j(\lambda^{(1)}) < c_j(\lambda^{(2)})$ for at least one $j \in \{1, 2, \dots, m\}$
- A candidate λ^* that is not dominated by any other candidate is called **Pareto optimal**.
- The set of all Pareto optimal candidates is called **Pareto set**
 $\mathcal{P} := \{\lambda \in \Lambda \mid \nexists \tilde{\lambda} \text{ with } c(\tilde{\lambda}) \prec c(\lambda)\}$
- $\mathcal{F} = c(\mathcal{P}) = \{c(\lambda) \mid \lambda \in \mathcal{P}\}$ is called **Pareto front**.

How to define optimality? I

Let $c = (\text{price}, -\text{rating})$. For some cases it is *clear* which point is the better one:

- The candidate $c^{(1)} = (89, -4.0)$ dominates $c^{(2)} = (108, -4.0)$: $c^{(1)}$ is not worse in any dimension and is better in one dimension. Therefore, $c^{(2)}$ gets **dominated** by $c^{(1)}$

$$c^{(2)} \prec c^{(1)}.$$




How to define optimality? II

For the points $c^{(1)} = (89, -4.0)$ and $c^{(2)} = (95, -4.5)$ we cannot say which one is better.


- We define the points as **equivalent** and write

$$c^{(1)} \not\prec c^{(2)} \text{ and } c^{(2)} \not\prec c^{(1)}.$$



How to define optimality? III

- The set of all equivalent points that are not dominated by another point is called the **Pareto front**.




Example: One objective function

We consider the minimization problem

$$\min_{\lambda} c(\lambda) = (\lambda - 1)^2, \quad 0 \leq \lambda \leq 3.$$


The optimum is at $\lambda^* = 1$.



Example: Two target functions I


We extend the above problem to two objective functions $c_1(\lambda) = (\lambda - 1)^2$ and $c_2(\lambda) = 3(\lambda - 2)^2$, thus

$$\min_{\lambda} c(\lambda) = (c_1(\lambda), c_2(\lambda)), \quad 0 \leq \lambda \leq 3.$$



Example: Two target functions II

We consider the functions in the objective function space $c(\Lambda)$ by drawing the objective function values $(c_1(\lambda), c_2(\lambda))$ for all $0 \leq \lambda \leq 3$.



The Pareto front is shown in green. The Pareto front cannot be *left* without getting worse in at least one objective function.

A-priori vs. A-posteriori

- The Pareto set is a set of equally optimal solutions.
- In many applications one is often interested in a **single** optimal solution.
- Without further information no unambiguous optimal solution can be determined.
→ The decision must be based on other criteria.

There are two possible approaches:

- **A-priori approach:** User preferences are considered **before** the optimization process
- **A-posteriori approach:** User preferences are considered **after** the optimization process


A-priori procedure I

Example: Weighted total

Prior knowledge: One rating point is worth 50 Euro to a customer.

→ We optimize the weighted sum:

$$\min_{\text{Hotel}} (\text{Price} / \text{Night}) - 50 \cdot \text{Rating}$$





Alternative a weighted sum: $\min_{\lambda \in \Lambda} \sum_{i=1}^m w_i c_i(\lambda)$ with $w_i \geq 0$

A-priori procedure II

Example: Lexicographic method

Prior knowledge: Customer prioritizes rating over price.

→ Optimize target functions one after the other.



A-priori procedure III

A-priori approach: Lexicographic method

$$\begin{aligned}c_1^* &= \min_{\lambda \in \Lambda} c_1(\lambda) \\c_2^* &= \min_{\lambda \in \{\lambda \mid c_1(\lambda) = c_1^*\}} c_2(\lambda) \\c_3^* &= \min_{\lambda \in \{\lambda \mid c_1(\lambda) = c_1^* \wedge c_2(\lambda) = c_2^*\}} c_3(\lambda) \\&\vdots\end{aligned}$$

But: Different sequences provide different solutions.

A-priori procedure IV

Summary a-priori approach:

- Implicit assumption: Single-objective optimization is *easy*.
- Only one solution is obtained, which depends on a-priori weights, order, etc.
- Several solutions can be obtained if weights, order, etc. are systematically varied.
- Usually not all non-dominated candidates can be found by these methods.

A-posteriori procedure I



A-posteriori methods try to

- find the set of **all** optimal candidates (the Pareto set),
- select (if necessary) an optimal candidate based on prior knowledge or individual preferences.
- Implicit assumption: Specifying your hidden preferences / making a selection from a pool of candidates is easier, if you see the non-dominated solutions.

A-posteriori methods are therefore the more generic approach to solving a multi-criteria optimization problem.

A-posteriori procedure II

Example: A user is displayed all Pareto optimal hotels (left) and chooses an optimal candidate (right) based on his hidden preferences or additional criteria (e.g. location of the hotel).




Evaluation of solutions I

A common metric for evaluating the performance of a set of candidates $\mathcal{P} \subset \mathbf{\Lambda}$ is the **dominated hypervolume**

$$S(\mathcal{P}, R) = \Lambda \left(\bigcup_{\tilde{\lambda} \in \mathcal{P}} \{ \lambda | \tilde{\lambda} \prec \lambda \prec R \} \right),$$

where Λ is the Lebesgue measure.



○ Reference point

Evaluation of solutions II

- HV is calculated w.r.t the reference point R , which often reflects in each component the natural maximum of the respective objective – if possible
- The dominated hypervolume is also often called **S-Metric**.
- Computation of HV scales exponentially in the number of objective functions $\mathcal{O}(n^{m-1})$.
- Fast approximations exist for small values of m and especially for machine learning applications we rarely optimize $m > 3$ objectives.





Multi-criteria Optimization

Evolutionary Approaches

Bernd Bischl Frank Hutter Lars Kotthoff
Marius Lindauer Joaquin Vanschoren

A-posteriori methods and evolutionary algorithms I

Evolutionary multi-objective algorithms (EMOAs) evolve a diverse population over time to approximate the Pareto front.



Algorithm 1 Basic EA template loop

- 1: Init and eval population $\mathcal{P}_0 \subset \Lambda$ with $|\mathcal{P}| = \mu$
 - 2: $t \leftarrow 0$
 - 3: **repeat**
 - 4: Select parents and generate offspring \mathcal{Q}_t with $|\mathcal{Q}_t| = \lambda$
 - 5: Select μ survivors \mathcal{P}_{t+1}
 - 6: $t \leftarrow t + 1$
 - 7: **until** Stop criterion fulfilled
-

- Note that (as in the EA lecture unit) we are using somewhat non-standard notation here.
- Nearly all steps in the above template work also for EMOAs but both parent and survival selection are now less obvious. How do we rank under multiple objectives?


The **non-dominated sorting genetic algorithm (NSGA-II)** was published by [Dep et al. 2002].

- Follows a $(\mu + \lambda)$ strategy.
- All previously discussed variation strategies can be used; the original paper uses tournament selection, polynomial mutation and simulated binary crossover.
- Parent and survival selection rank candidates by
 - 1 **Non-dominated sorting** as main criterion
 - 2 **Crowding distance assignment** as tie breaker

NSGA-II: non-dominated sorting I


NDS partitions an objective space set into fronts $\mathcal{F}_1 \prec \mathcal{F}_2 \prec \mathcal{F}_3 \prec \dots$

- \mathcal{F}_1 is non-dominated, each $\lambda \in \mathcal{F}_2$ is dominated, but only by points in \mathcal{F}_1 , each $\lambda \in \mathcal{F}_3$ is dominated, but only by points in \mathcal{F}_1 and \mathcal{F}_2 , and so on.
- We can easily compute the partitioning by computing all non-dominated points \mathcal{F}_1 , removing them, then computing the next layer of non-dominated points \mathcal{F}_2 , and so on.



NSGA-II: non-dominated sorting II

How does survival selection now work? We fill μ places one by one with $\mathcal{F}_1, \mathcal{F}_2, \dots$ until a front can no longer **fully** survive (here: \mathcal{F}_3).




Which individuals survive from \mathcal{F}_3 ? \rightarrow **crowding sort**

NB: the same principle to rank individuals is applied in tournament selection in parent selection.

NSGA-II: crowding distance I

Idea: Add *good* representatives of front \mathcal{F}_3 , define this as points of "low density" in c-space.




Left: Not good, points very close together. Right: better.

NSGA-II: crowding distance II

For each objective c_j

- Sort points by c_j
- Normalize scores to $[0,1]$
- Assign border points (which have score 0 or 1) a CD of ∞ (they should always be selected, if possible)
- Each point gets a distance score, which is the distance between its 2 next-neighbors w.r.t. the sorting of c_j

For each point, all of its m distance scores are summed up (or averaged) and points are ranked w.r.t. to this overall score.



Red: Point with high CD. Blue: Low CD.

Selection criteria: contribution to the hypervolume I


SMS-EMOA

(S-Metric-Selection-EMOA) [Beume et al. 2007]

is a $(\mu + 1)$ EMOA and evaluates fitness of an individual $\lambda \in \mathcal{P} \subset \Lambda$ based on its contribution to the dominated HV:

$$\Delta s(\lambda, \mathcal{P}) = S(\mathcal{P}, R) - S(\mathcal{P} \setminus \{\lambda\}, R).$$

- Dark rectangles: HV contribution of dots.
- Grey point: reference point.
- The HVC contribution is the volume of space that is dominated only by λ , and nothing else.
- $\tilde{\lambda}$ has lowest S-metric contribution.



SMS-EMOA algorithm I

Algorithm 2 SMS-EMOA

- 1: Generate start population \mathcal{P}_0 of size μ
 - 2: $t \leftarrow 0$
 - 3: **repeat**
 - 4: Generate **one** individual \mathbf{q} by recombination and mutation of \mathcal{P}_t
 - 5: $\{\mathcal{F}_1, \dots, \mathcal{F}_k\} \leftarrow \text{NDS}(\mathcal{P}_t \cup \{\mathbf{q}\})$
 - 6: $\tilde{\lambda} \leftarrow \operatorname{argmin}_{\lambda \in \mathcal{F}_k} \Delta s(\lambda, \mathcal{F}_k)$
 - 7: $\mathcal{P}_{t+1} \leftarrow (\mathcal{P}_t \cup \{\mathbf{q}\}) \setminus \{\tilde{\lambda}\}$
 - 8: $t \leftarrow t + 1$
 - 9: **until** Termination criterion fulfilled
-

- L5: the set of temporary $(\mu + 1)$ individuals is partitioned by NDS into k fronts $\mathcal{F}_1, \dots, \mathcal{F}_k$.
- L6-7: In last front, find $\tilde{\lambda} \in \mathcal{F}_k$ with smallest HV contribution - and kill it.
- Fitness of an individual is mainly the rank of its front and HV contribution as tie-breaker.

Multi-criteria Optimization

Practical Applications

Bernd Bischl Frank Hutter Lars Kotthoff
Marius Lindauer Joaquin Vanschoren

Practical Applications in Machine Learning I

ROC Optimization: Balance *true positive* and *false positive* rates

- Typically unbalanced classification tasks with unspecified costs.
- Could also use other ROC metrics, e.g., *positive predicted value* or *false discovery rate*.

Efficient Models: Balance *predictive performance* with *prediction time*, *energy consumption* and/or *model size*.

- Time: Models in production need to predict fast.
- Size / Energy consumption: Models should be deployed on a mobile/edge device and not use much power.

Sparse Models: Balance *predictive performance* and *number of used features*, either for cost efficiency, but often also for interpretability.

Fair Models: Balance *predictive performance* and *fairness*.

- Model has to be fair regarding subgroups in the data, e.g. gender.
- Many different approaches to quantify fairness exist.

ROC Optimization - Setup

Again, we want to train a *spam detector* on the popular Spam dataset¹.

- Learning algorithm: SVM with RBF kernel.
- Hyperparameters to optimize:
 - cost $[2^{-15}, 2^{15}]$
 - γ $[2^{-15}, 2^{15}]$
 - Threshold t $[0, 1]$
- Objective: *minimize* false positive rate (FPR) and *maximize* true positive rate (TPR), evaluated through 5-fold CV
- Optimizer: Multi-criteria Bayesian optimization:
 - ▶ ParEGO with $\rho = 0.05$, $s = 100000$.
 - ▶ Acquisition function u : *Confidence Bound* with $\alpha = 2$.
 - ▶ Budget: 100 evaluations
- Tuning is conducted on a training holdout and all hyperparameter configurations on the estimated Pareto front are validated on an outer validation set.

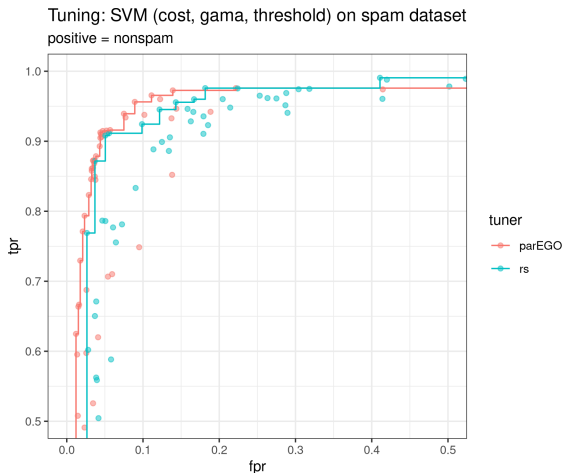
The threshold t could be separately optimized post-hoc.

¹<https://archive.ics.uci.edu/ml/datasets/spambase>

ROC Optimization - Result I

We notice:


- Compared to *random search*: Many *ParEGO* evaluations are on the Pareto front.
- The Pareto front of *ParEGO* dominates most points from the *random search*.
- The dominated hypervolume to the reference point (0, 1) is:
 ParEGO: 0.965
 random search: 0.959



ROC Optimization - Result II

We validate the configurations on the estimated Pareto front on a holdout:


- The performance on the validation set varies slightly.
- The TPR got slightly better but the FPR got slightly worse.
- On the validation set, some configurations get dominated by others.



ROC Optimization - Result II

We validate the configurations on the estimated Pareto front on a holdout:

- The performance on the validation set varies slightly.
- The TPR got slightly better but the FPR got slightly worse.
- On the validation set, some configurations get dominated by others.
- The dominated hypervolume of the validation set is:
ParEGO: 0.960
random search: 0.961



Efficient Models - Overview

- "Efficiency" can be:
 - ▶ Memory consumption of the model
 - ▶ Training or prediction time
 - ▶ Number of features needed
 - ▶ Energy consumption for prediction
 - ▶ ...
- Some hyperparameters have a strong impact on the efficiency of a model, e.g.,
 - ▶ Number of trees in *random forests* or *gradient tree boosting*,
 - ▶ Number, size and type of layers in *neural networks*,
 - ▶ L1 regularization penalties,
 - ▶ ...
- Other hyperparameters might have no influence on efficiency.
- Typical scenario: Optimize jointly over multiple algorithms of varying efficiency.


Efficient Models - Example: Feature Selection I

Goal of *feature selection*: Identify an informative feature subset with only a small drop in predictive performance compared to all features.

Find optimal hyperparameter setting λ and minimal feature subset s

$$\min_{\lambda \in \Lambda, s \in \{0,1\}^p} \left(\widehat{GE}(\mathcal{I}(\mathcal{D}, \lambda, s)), \frac{1}{p} \sum_{i=1}^p s_i \right)$$

- Problem: Feature selection and hyperparameter tuning are usually two separate steps.
- Solution: Identify an informative subset of features and a good hyperparameter configuration **simultaneously**.



Efficient Models - Example: Feature Selection II

Idea: *Multi-Objective Hyperparameter Tuning and Feature Selection using Filter Ensembles* [Binder et al. 2020]:


- Pre-calculate multiple ranked *feature filter values* .

RF Feature Importance: (2 3 4 1 5) x

AUC: (2 1 3 4 5) x

Information Gain: (1 3 5 2 4) x

1.9 2.6 3.9 1.7 4.9



0.7
0.2
0.1

\mathbf{w} : search space component

$$EF_j(\mathbf{w}) = \sum_{m=1}^M w_m F^m(\mathcal{D})_j.$$


- New joint hyperparameter vector: $\boldsymbol{\lambda} = (\tilde{\boldsymbol{\lambda}}, w_1, \dots, w_p, \tau)$
 - ▶ Hyperparameters of learner: $\tilde{\boldsymbol{\lambda}}$
 - ▶ Weight of each *feature filter value* vector: (w_1, \dots, w_p)
 - ▶ Fraction of features to keep τ

Efficient Models - Example: Feature Selection III

Combined feature selection and hyperparameter optimization on Sonar dataset².

- Learning algorithm: SVM with RBF kernel.
- Hyperparameters to optimize:

cost	$[2^{-10}, 2^{10}]$
γ	$[2^{-10}, 2^{10}]$
(w_1, \dots, w_p)	$[0, 1]^p$
τ	$[0, 1]^p$
- Objective: minimize *misclassification* and *fraction of features selected*
- Optimizer: *ParEGO* with *random forest* surrogate, LCB acquisition function, 15 batch proposals, budget: 2000 evaluations



²Only the tuning error is shown here

Efficient Models - Example: FLOPS

Goal: Optimize prediction accuracy and number of floating point operations (FLOPs) [Wang et al. 2019].

Data: Image Classification on CIFAR-10.


Learner: *DenseNet* - Densely Connected Convolutional Network [Huang et al. 2018].

- Composed out of 4 *dense blocks*.
- A dense block consists of multiple convolutional layers where the inputs for each layer are all feature maps of all preceding layers in the block.
- Dense blocks are connected via convolutional and max pooling layer.

Training: 300 Epochs with a batch size of 128 and initial learning rate of 0.1.

Efficient Models - Example: FLOPS

- Objective: *accuracy* vs. *FLOPS* (floating point operations, per observation)
- Search Space:
 - growth rate (k) [8, 32]
 - layers in first block [4, 6]
 - layers in second block [4, 12]
 - layers in third block [4, 24]
 - layers in fourth block [4, 16]
- Tuner: *Particle Swarm Optimization* with a population size of 20 and 400 evaluations.






The growth rate is the number of output feature maps in each layer of a block

Fair Models - The Adult dataset

Dataset: Adult

- Source: US Census database, 1994, <https://www.openml.org/d/1590>.
- 48842 observations
- Target: binary, income above 50k
- 14 features: age, education, hours.per.week, marital.status, native.country, occupation, race, relationship, sex, ...



Fair Models - Setup I

A fair model for income prediction on binarized target.

- Learner: *eXtreme Gradient Boosting*

- Hyperparameters to optimize:

eta	$[0.01, 0.2]$
gamma	$[2^{-7}, 2^6]$
max_depth	$\{2, \dots, 20\}$
colsample_bytree	$[0.5, 1]$
colsample_bylevel	$[0.5, 1]$
lambda	$[2^{-10}, 2^{10}]$
alpha	$[2^{-10}, 2^{10}]$
subsample	$[0.5, 1]$

- Objective: minimize *misclassification error* and *unfairness*

Fair Models - Setup II

- Careful: Usually this data would be used to model the relation between person characteristics and income, then to **discuss and study** by careful inference - to **figure out if** something like e.g. a "gender pay gap" exists.
- Here, in our toy example we **pretend** now that we would like to create a automatic "assignment algorithm" for salary - maybe not totally unrealistic nowadays? In **such** a scenario, biasing the prediction by **incorporating fairness** might be of interest.
- Here, a simplified proxy for fairness is defined as the absolute difference in F1-Scores between female (f) and male (m) population (low is good):

$$L_{\text{fair}} := |L_{\text{F1}}(y_f, \hat{f}(\mathbf{x}_f)) - L_{\text{F1}}(y_m, \hat{f}(\mathbf{x}_m))|$$

Fair Models - Results




Figure: Pareto fronts after 20, 70 and 120 tuning iterations.

- Optimizer: ParEGO with random forest surrogate and restricted range of projections to $[0.1, 0.9]$ (No interest in very unfair or bad configurations).
- Here, the hyperparameters actually have an effect on the defined *fairness measure*.
- However, this is often not the case or not enough to ensure a fair model.