

# AutoML: Evaluation

## Benchmarking and Comparing Learners

Bernd Bischl   Frank Hutter   Lars Kotthoff  
Marius Lindauer   Joaquin Vanschoren

# Benchmark Experiments

- different learning algorithms applied to one or more data sets to compare and rank their performances
- synchronized train and test sets, i.e. the same resampling method with the same train-test splits should be used to determine performance

**Example:** Benchmark results (per CV-fold) of CART and random forest using 2-fold CV with MSE as performance measure:

data set	k-th fold	MSE (rpart)	MSE (randomForest)
BostonHousing	1	29.4	17.13
BostonHousing	2	20.5	8.90
mtcars	1	35.0	7.53
mtcars	2	38.9	6.73

# Hypothesis Testing in Benchmarking I

We want to know if the difference in performance between models (or algorithms) is significant or only by chance.

**Null Hypothesis Statistical Testing (NHST)** in Benchmarking:

- formulate null hypothesis  $H_0$  (e.g. the expected generalization error of two algorithms is equivalent) and alternative hypothesis  $H_1$
- use hypothesis test to reject  $H_0$  (not rejecting  $H_0$  does not mean that we accept it)
- rejecting  $H_0$  gives some confidence that the observed results may not be random

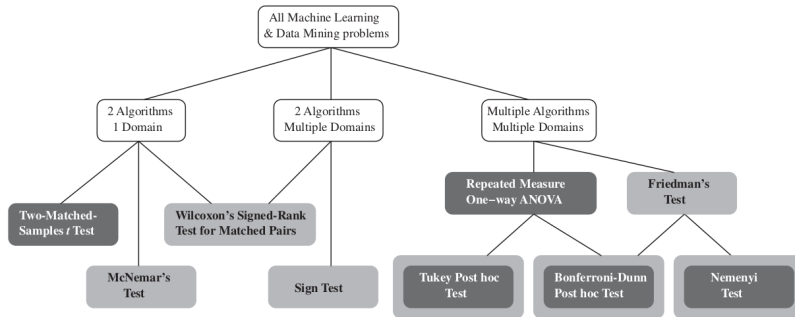
Typical example in machine learning:

- $H_0$ : on average, model 1 does not perform better than model 2
- $H_1$ : on average, model 1 outperforms model 2
- Aim: Reject  $H_0$  with confidence level of  $1 - \alpha$  (common values for  $\alpha$  include 0.05 and 0.01)

# Hypothesis Testing in Benchmarking II

Selection of an appropriate hypothesis test is at least based on the type of problem, i.e. if the aim is to compare

- 2 models / algorithms on a single domain (i.e. on a single data set)
- 2 algorithms across different domains (i.e. on multiple data sets)
- multiple algorithms across different domains / data sets



Legend:

Parametric Test

Parametric and Nonparametric

Nonparametric

# McNemar Test I

- non-parametric test used on paired dichotomous nominal data; does not make any distributional assumptions beyond statistical independence of samples
- pairs are e.g. labels predicted by different models on the same data
- compares the classification accuracy of two **models**
- both models trained and evaluated on the exact same training and test set; **contingency table** based on two paired vectors that indicate whether each model predicted an observation correctly

	Model 2 correct	Model 2 wrong
Model 1 correct	A	B
Model 1 wrong	C	D

- A: #obs. correctly classified by both
- B: #obs. only correctly classified by model 1
- C: #obs. only correctly classified by model 2
- D: #obs. misclassified by both


# McNemar Test II

Error of each model can be computed as follows:

- Model 1:  $(C+D)/(A+B+C+D)$
- Model 2:  $(B+D)/(A+B+C+D)$

Even if the models have the **same** errors (indicating equal performance), cells B and C may be different because the models may misclassify different instances.

	Model 2 correct	Model 2 wrong
Model 1 correct	A	B
Model 1 wrong	C	D

 This work by Sebastian Raschke is licensed under a Creative Commons Attribution 4.0 International License.

McNemar tests the following hypothesis:

- $H_0$  : both models have the same performance (we expect  $B = C$ )
- $H_1$  : performances of the two models are not equal

The test statistic is computed as

$$\chi_{Mc}^2 = \frac{(|B-C|-1)^2}{B+C} \sim \chi_1^2.$$

**Note:** The McNemar test should only be used if  $B + C > 20$ .

# McNemar Test III

## Example:

		Random Forest	
		0	1
Tree	0	30	5
	1	17	42

Calculating the test statistic:

$$\chi_{Mc}^2 = \frac{(|5 - 17| - 1)^2}{5 + 17} = 5.5 > 3.841 = \chi_{1,0.95}^2$$

We can reject  $H_0$  at a significance level of 0.05, i.e. we reject the hypothesis that the tree and the random forest have the same performance.

Significance level must be chosen before applying the test (avoid p-value hacking).

# Two-Matched-Samples t-Test I

- two-matched-samples t-test (i.e. a paired t-test) is the simplest hypothesis test to compare two **models** on a single test set based on arbitrary performance measures
- parametric test and distributional assumptions must be made (which are often problematic):
  - (pseudo-)normality usually met when sample size  $> 30$
  - i.i.d. samples usually met if loss of individual observations from single test set considered
  - equal variances of populations can be investigated through plots



## Two-Matched-Samples t-Test II

Compare two different models  $\hat{f}_1$  and  $\hat{f}_2$  w.r.t. performance measure calculated on test set of size  $n_{\text{test}}$ :

- $H_0: GE(\hat{f}_1) = GE(\hat{f}_2)$  vs.  $H_1: GE(\hat{f}_1) \neq GE(\hat{f}_2)$
- test statistic  $T = \sqrt{n_{\text{test}}} \frac{\bar{d}}{\sigma_d}$  where
  - ▶ mean performance difference of both models is  $\bar{d} = \hat{GE}_{\mathcal{D}_{\text{test}}}(\hat{f}_1) - \hat{GE}_{\mathcal{D}_{\text{test}}}(\hat{f}_2)$
  - ▶ standard deviation of this mean difference is

$$\sigma_d = \sqrt{\frac{1}{n_{\text{test}} - 1} \sum_{i=1}^{n_{\text{test}}} (d_i - \bar{d})^2},$$

where  $d_i = L(y^{(i)}, \hat{f}_1(\mathbf{x}^{(i)})) - L(y^{(i)}, \hat{f}_2(\mathbf{x}^{(i)}))$  and  $\bar{d} = \sum_{i=1}^{n_{\text{test}}} d_i$

**Note:**  $d_i$  is the difference of the outer loss of individual observations from the test set between the two models to be compared.

## Two-Matched-Samples t-Test III

- could also use a  **$k$ -fold CV paired t-test** to compare two **algorithms** (instead of two models) on single data set
- instead of comparing outer loss of individual observations, compare generalization errors per CV fold (i.e.  $k$  generalization errors for  $k$  CV folds)
- performance differences are not independent across CV folds due to overlapping training sets (which violates the assumption of i.i.d. samples)
- to partly overcome issue of overlapping training sets across folds, Dietterich<sup>1</sup> suggests using 5 times 2-fold CV so that at least within each repetition neither training nor test sets overlap

---

<sup>1</sup>Dietterich (1998). Approximate statistical tests for comparing supervised classification learning algorithms.

# Friedman Test I

Compare multiple classifiers on multiple data sets:

- $H_0$  : all algorithms are equivalent in their performance and hence their average ranks should be equal
- $H_1$  : the average ranks for at least one algorithm is different

To evaluate  $n$  data sets and  $k$  algorithms:

- rank each algorithm on each data set from best-performing algorithm (rank 1) to worst-performing algorithm using any performance measure
- $R_{ij}$  is the rank of algorithm  $j$  on data set  $i$
- if there is a  $d$ -way tie after rank  $r$ , assign rank of  $[(r + 1) + (r + 2) + \dots + (r + d)] / d$  to each tied classifier

## Friedman Test II

Can now compute:

- overall mean rank  $\bar{R} = \frac{1}{nk} \sum_{i=1}^n \sum_{j=1}^k R_{ij}$
- sum of squares total  $SS_{Total} = n \sum_{j=1}^k (\bar{R}_{.j} - \bar{R})^2$  where  $\bar{R}_{.j} = \frac{1}{n} \sum_{i=1}^n R_{ij}$
- sum of squares error  $SS_{Error} = \frac{1}{n(k-1)} \sum_{i=1}^n \sum_{j=1}^k (R_{ij} - \bar{R})^2$

Test statistic calculated as:

$$\chi_F^2 = \frac{SS_{Total}}{SS_{Error}} \sim \chi_{k-1}^2 \text{ for large } n (>15) \text{ and } k (>5)$$

For smaller  $n$  and  $k$ , the  $\chi^2$  approximation is imprecise and a look up of  $\chi_F^2$  values that were approximated specifically for the Friedman test is suggested.

# Post-Hoc Tests I

- Friedman test checks if all algorithms are ranked equally
- does not allow to identify best-performing algorithm

→ post-hoc tests

## Post-hoc Nemenyi test:

- compares all pairs of algorithms to find best-performing algorithm after  $H_0$  of the Friedman-test was rejected
- for  $n$  data sets and  $k$  algorithms,  $\frac{k(k-1)}{2}$  comparisons
- calculate average rank of algorithm  $j$  on all  $n$  data sets:  $\bar{R}_{.j} = \frac{1}{n} \sum_{i=1}^n R_{ij}$

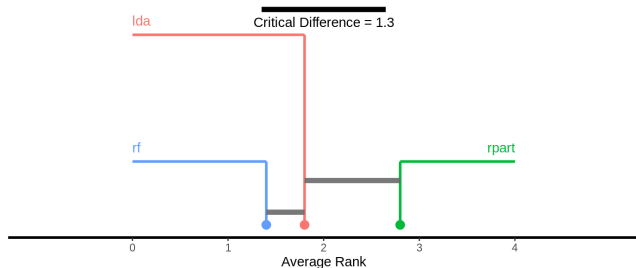
For any two algorithms  $j_1$  and  $j_2$ , test statistic computed as:

$$q = \frac{\bar{R}_{.j_1} - \bar{R}_{.j_2}}{\sqrt{\frac{k(k+1)}{6n}}}$$

# Post-Hoc Tests II

## Critical Difference Plot:

- quick way to see what differences are significant across all compared learners
- all learners that do not differ by at least the critical difference are connected by line
- a learner not connected to another learner and of lower rank can be considered better according to the chosen significance level



# Post-Hoc Tests III

## Post-hoc Bonferonni-Dunn test:

- compares all algorithms with baseline (i.e.  $k - 1$  comparisons)
- used after Friedman test to find which algorithms differ from the baseline significantly
- uses Bonferonni correction to prevent randomly accepting one of the algorithms as significant due to multiple testing


The test statistic is the same as before:

$$q = \frac{\bar{R}_{.j_1} - \bar{R}_{.j_2}}{\sqrt{\frac{k(k+1)}{6n}}}.$$

The performance of  $j_1$  and  $j_2$  are significantly different when  $|q| > q_\alpha$ , where the critical value  $q_\alpha$  is obtained from a table of the studentized range statistic divided by  $\sqrt{2}$ .

# Post-Hoc Tests IV

- learners within the baseline interval (gray line) perform not significantly different from the baseline

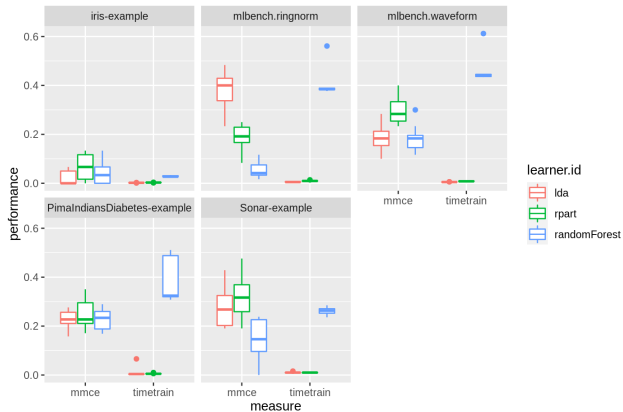




# Comparing Visually I

It can be helpful to inspect distributions visually for additional insights, e.g.

## Boxplots



# Comparing Visually II

## Rank plots

