

# AutoML: Hyperparameter Optimization

## Evolutionary Algorithms

Bernd Bischl   Frank Hutter   Lars Kotthoff  
Marius Lindauer   Joaquin Vanschoren

# Evolutionary algorithms

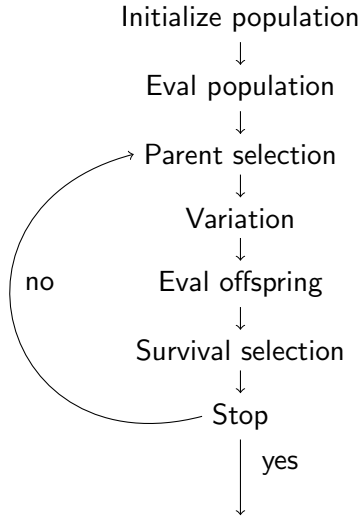
**Evolutionary algorithms** (EA) are a class of stochastic, metaheuristic optimization techniques whose mode of operation is inspired by the evolution of natural organisms.

History of evolutionary algorithms:

- **Genetic algorithms:** Use binary problem representation, therefore closest to the biological model of evolution.
- **Evolution strategies:** Use direct problem representation, e.g., vector of real numbers.
- **Genetic programming:** Create structures that convert an input into a fixed output (e.g. computer programs); solution candidates are represented as trees.
- **Evolutionary programming:** Similar to GP, but solution candidates are not represented by trees, but by finite state machines.

The boundaries between the terms become increasingly blurred and are often used synonymously.

# Structure of an evolutionary algorithm



# Notation and Terminology

Symbols	EA Terminology
Solution candidate $\lambda \in \Lambda$	Chromosome of an individual
$\lambda_i$	$i$ -th gene of chromosome
Set of candidates $\mathcal{P}$ with $\mu =  \mathcal{P} $	Population and size
$\lambda$	Number of generated offsprings
$c : \Lambda \rightarrow \mathbb{R}$	Fitness function

$$c(\lambda) = \widehat{GE}_{\mathcal{D}_{\text{test}}}(\mathcal{I}(\mathcal{D}_{\text{train}}, \lambda))$$

Notation clash:

- In EAs the objective function is often denoted  $f(x)$ .
- As these symbols are used for ML already we use  $c(\lambda)$  and  $\lambda$  instead of  $f$  and  $x$ .
- Be careful: The offspring size  $\lambda$  is different from the candidate  $\lambda$  (bold symbol!).

## Step 1: Initialize population


- A evolutionary algorithm is started by generating a initial population  $\mathcal{P} = \{\boldsymbol{\lambda}^{(1)}, \dots, \boldsymbol{\lambda}^{(\mu)}\}$ .
- Usually we sample this uniformly at random.
- We could introduce problem prior knowledge via a smarter init procedure.
- This population is evaluated, i.e., the objective function is computed for every individual in the initial population.
- The initialization can have a large influence on the quality of the found solution, so many EAs employ *restarts* with new randomly generated populations.

## Step 2: Parent selection I

In the first step of an iteration,  $\lambda$  parents are chosen, who create offspring in the next step.

Possibilities for selection of parents:

- **Neutral selection:** choose individual with a probability  $1/\mu$ .
- **Fitness-proportional selection:** draw individuals with probability proportional to their fitness.
- **Tournament Selection:** randomly select  $k$  individuals for a "Tournament Group". Of the drawn individuals, the best one (with the highest fitness value) is then chosen. Procedure is performed  $\lambda$ -times.



## Step 3: Variation

New individuals are now generated from the parent population. This is done by

- Recombination/Crossover: combine two parents into one offspring.
- Mutation: (locally) change an individual.

Sometimes only one operation is performed.

# Recombination for numeric representations

Two individuals  $\lambda, \tilde{\lambda} \in \mathbb{R}^n$  in numerical representation can be recombined as follows:

- **Uniform crossover**: choose gene  $i$  with probability  $p$  of 1st parent and probability  $1 - p$  of 2nd parent.
- **Intermediate recombination**: new individual is created from the mean value of two parents  $\frac{1}{2}(\lambda + \tilde{\lambda})$
- **Simulated Binary Crossover (SBX)**: generate **two offspring** based on

$$\bar{\lambda} \pm \frac{1}{2}\beta(\tilde{\lambda} - \lambda)$$


with  $\bar{\lambda} = \frac{1}{2}(\lambda + \tilde{\lambda})$  and  $\beta$  randomly sampled from a certain distribution.



# Mutation for numeric representations

**Mutation:** individuals are changed, for example for  $\lambda \in \mathbb{R}^n$

- **Uniform mutation:** choose a random gene  $\lambda_i$  and replace it with a value uniformly distributed (within the feasible range).
- **Gauss mutation:**  $\tilde{\lambda} = \lambda \pm \sigma \mathcal{N}(0, \mathbf{I})$
- **Polynomial mutation:** polynomial distribution instead of normal distribution



Source: K. Deb, Analysing mutation schemes for real-parameter genetic algorithms, 2014

# Recombination for bit strings

Two individuals  $\lambda, \tilde{\lambda} \in \{0, 1\}^n$  encoded as bit strings can be recombined as follows:

- **1-point crossover:** select crossover  $k \in \{1, \dots, n-1\}$  randomly and select the first  $k$  bits from 1st parent, the last  $n-k$  bits from 2nd parent.

$$\begin{array}{ccc} 1 & 1 & 1 \\ 0 & 0 & 0 \\ \hline 0 & 1 & \Rightarrow 1 \\ 1 & 1 & 1 \\ 1 & 0 & 0 \end{array}$$

- **Uniform crossover:** select bit  $i$  with probability  $p$  of 1st parent and  $1-p$  of 2nd parent.

$$\begin{array}{ccc} 1 & 0 & 1 \\ 0 & 0 & 0 \\ 0 & 1 & \Rightarrow 1 \\ 0 & 1 & 1 \\ 1 & 0 & 1 \end{array}$$

# Mutation for bit strings

An individual  $\lambda \in \{0, 1\}^n$  encoded as a bit string can be mutated as follows:

- **Bitflip:** for each index  $k \in \{1, \dots, n\}$ : bit  $k$  is flipped with probability  $p \in (0, 1)$ .

1		0
0		0
0	$\Rightarrow$	0
0		1
1		1

## Step 4: Survival selection


Now individuals are chosen who survive. Two common strategies are:

- $(\mu, \lambda)$ -**selection**: we select from the  $\lambda$  descendants the  $\mu$  best ( $\lambda \geq \mu$  necessary). **But:** best overall individual can get lost!
- $(\mu + \lambda)$ -**selection**:  $\mu$  parents and  $\lambda$  offspring are lumped together and the  $\mu$  best individuals are chosen. Best individual safely survives.

# Example of an evolutionary algorithm I


In the following, a (simple) EA is shown on the 1-dim Ackley function, optimized on  $[-30, 30]$ .

Usually for the optimization of a function  $c : \mathbb{R}^n \rightarrow \mathbb{R}$  individuals are coded as real vectors  $\lambda \in \mathbb{R}^n$ , so here we use simply one real number as chromosome.




# Example of an evolutionary algorithm II

Randomly init population with size  $\mu = 20$ .




# Example of an evolutionary algorithm III

We choose  $\lambda = 5$  offspring by neutral selection (red individuals).



# Example of an evolutionary algorithm IV


We use a Gauss mutation with  $\sigma = 2$  and do not apply a recombination.





# Example of an evolutionary algorithm V

We use a  $(\mu + \lambda)$  selection. The selected individuals are green.



# Evolutionary Algorithms

## Advantages

- Conceptually simple, yet powerful enough to solve complex problems (including HPO)
- All parameter types possible in general
- Highly parallelizable (depends on  $\mu$ )
- Allows customization via specific variation operators

## Disadvantages

- Less theory available (for realistic, complex EAs)
- Can be hard to get balance between exploration and exploitation right
- Can have quite a few control parameters, hard to set them correctly
- Customization necessary for complex problems
- Not perfectly suited for expensive problems like HPO, as quite a higher number of evaluations is usually needed for appropriate convergence / progress