


# AutoML: Neural Architecture Search (NAS)

DARTS: Differentiable Architecture Search


Bernd Bischl   Frank Hutter   Lars Kotthoff  
Marius Lindauer   Joaquin Vanschoren

# DARTS: Differentiable Architecture Search [Liu et al, 2018]

- Use one-shot model with continuous architecture weight  $\alpha$  for each operator



(a) Initialization




(b) Search end


# DARTS: Differentiable Architecture Search [Liu et al, 2018]

- Use one-shot model with continuous architecture weight  $\alpha$  for each operator

$$x^{(j)} = \sum_{i < j} \tilde{o}^{(i,j)}(x^{(i)}) = \sum_{i < j} \sum_{o \in \mathcal{O}} \frac{\exp(\alpha_o^{(i,j)})}{\sum_{o' \in \mathcal{O}} \exp(\alpha_{o'}^{(i,j)})} o(x^{(i)})$$



(a) Initialization




(b) Search end


# DARTS: Differentiable Architecture Search [Liu et al, 2018]

- Use one-shot model with continuous architecture weight  $\alpha$  for each operator

$$x^{(j)} = \sum_{i < j} \tilde{o}^{(i,j)}(x^{(i)}) = \sum_{i < j} \sum_{o \in \mathcal{O}} \frac{\exp(\alpha_o^{(i,j)})}{\sum_{o' \in \mathcal{O}} \exp(\alpha_{o'}^{(i,j)})} o(x^{(i)})$$



(a) Initialization




(b) Search end

- By optimizing the architecture weights  $\alpha$ , DARTS assigns importance to each operation
  - ▶ Since the  $\alpha$  are continuous, we can optimize them with gradient descent


# DARTS: Differentiable Architecture Search [Liu et al, 2018]

- Use one-shot model with continuous architecture weight  $\alpha$  for each operator


$$x^{(j)} = \sum_{i < j} \tilde{o}^{(i,j)}(x^{(i)}) = \sum_{i < j} \sum_{o \in \mathcal{O}} \frac{\exp(\alpha_o^{(i,j)})}{\sum_{o' \in \mathcal{O}} \exp(\alpha_{o'}^{(i,j)})} o(x^{(i)})$$



(a) Initialization



(b) Search end



(c) Final cell

- By optimizing the architecture weights  $\alpha$ , DARTS assigns importance to each operation
  - ▶ Since the  $\alpha$  are continuous, we can optimize them with gradient descent
- In the end, DARTS discretizes to obtain a single architecture (c)

# DARTS: Architecture Optimization

- The optimization problem ( $a \rightarrow b$ ) is a bi-level optimization problem:

$$\begin{aligned} & \min_{\alpha} \mathcal{L}_{\text{val}}(w^*(\alpha), \alpha) \\ \text{s.t. } & w^*(\alpha) \in \operatorname{argmin}_w \mathcal{L}_{\text{train}}(w, \alpha) \end{aligned}$$

# DARTS: Architecture Optimization

- The optimization problem ( $a \rightarrow b$ ) is a bi-level optimization problem:

$$\begin{aligned} & \min_{\alpha} \mathcal{L}_{\text{val}}(w^*(\alpha), \alpha) \\ & \text{s.t. } w^*(\alpha) \in \operatorname{argmin}_w \mathcal{L}_{\text{train}}(w, \alpha) \end{aligned}$$

- This is solved using alternating SGD steps on architectural parameters  $\alpha$  and weights  $w$

---

**Algorithm:** DARTS 1st order

---

**while** *not converged* **do**

    Update one-shot weights  $\mathbf{w}$  by  $\nabla_{\mathbf{w}} \mathcal{L}_{\text{train}}(\mathbf{w}, \alpha)$

    Update architectural parameters  $\alpha$  by  $\nabla_{\alpha} \mathcal{L}_{\text{valid}}(\mathbf{w}, \alpha)$

**return**  $\arg \max_{o \in \mathcal{O}} \alpha_o^{(i,j)}$  for each edge  $(i, j)$

---

# DARTS: Architecture Optimization

- The optimization problem ( $a \rightarrow b$ ) is a bi-level optimization problem:

$$\begin{aligned} & \min_{\alpha} \mathcal{L}_{\text{val}}(w^*(\alpha), \alpha) \\ & \text{s.t. } w^*(\alpha) \in \operatorname{argmin}_w \mathcal{L}_{\text{train}}(w, \alpha) \end{aligned}$$

- This is solved using alternating SGD steps on architectural parameters  $\alpha$  and weights  $w$

---

**Algorithm:** DARTS 1st order

---

**while** *not converged* **do**

    Update one-shot weights  $\mathbf{w}$  by  $\nabla_{\mathbf{w}} \mathcal{L}_{\text{train}}(\mathbf{w}, \alpha)$

    Update architectural parameters  $\alpha$  by  $\nabla_{\alpha} \mathcal{L}_{\text{valid}}(\mathbf{w}, \alpha)$

**return**  $\arg \max_{o \in \mathcal{O}} \alpha_o^{(i,j)}$  for each edge  $(i, j)$

---

- Note: there is no theory showing that this process converges



# Strong performance on some benchmarks

- E.g., original CNN search space
  - ▶ 8 operations on each MixedOp
  - ▶ 28 MixedOps in total
  - ▶  $> 10^{23}$  possible architectures
- Performance
  - ▶  $< 3\%$  error on CIFAR-10 in less than 1 GPU day of search




Figure 4: Normal cell learned on CIFAR-10.




Figure 5: Reduction cell learned on CIFAR-10.

## Issues – Non-robust behaviour


- DARTS is very sensitive w.r.t. its own hyperparameters (e.g. one-shot learning rate,  $L_2$  regularization, etc.)

## Issues – Non-robust behaviour

- DARTS is very sensitive w.r.t. its own hyperparameters (e.g. one-shot learning rate,  $L_2$  regularization, etc.)
  - Tuning these hyperparameters for every new task/search space is computationally expensive


## Issues – Non-robust behaviour

- DARTS is very sensitive w.r.t. its own hyperparameters (e.g. one-shot learning rate,  $L_2$  regularization, etc.)
  - Tuning these hyperparameters for every new task/search space is computationally expensive
  - DARTS may return degenerate architectures, e.g., cells composed with only skip connections






# Issues – Non-robust behaviour

- DARTS is very sensitive w.r.t. its own hyperparameters (e.g. one-shot learning rate,  $L_2$  regularization, etc.)
  - Tuning these hyperparameters for every new task/search space is computationally expensive
  - DARTS may return degenerate architectures, e.g., cells composed with only skip connections
- RobustDARTS [Zela et al, 2020] – tracks the curvature of the validation loss and stops the search early based on that



# Issues – Non-robust behaviour

- DARTS is very sensitive w.r.t. its own hyperparameters (e.g. one-shot learning rate,  $L_2$  regularization, etc.)
  - Tuning these hyperparameters for every new task/search space is computationally expensive
  - DARTS may return degenerate architectures, e.g., cells composed with only skip connections
- RobustDARTS [Zela et al, 2020] – tracks the curvature of the validation loss and stops the search early based on that
- SmoothDARTS [Chen and Hsieh, 2020] – applies random perturbation and adversarial training to avoid bad regions



## Issues – Memory constraints

- DARTS keeps the **entire one-shot model in memory**, together with its computed tensors

## Issues – Memory constraints

- DARTS keeps the **entire one-shot model in memory**, together with its computed tensors
  - This constrains the search space size and the fidelity used to train the one-shot model
  - **Impossible** to run on large datasets as ImageNet



## Issues – Memory constraints

- DARTS keeps the **entire one-shot model in memory**, together with its computed tensors
  - This constrains the search space size and the fidelity used to train the one-shot model
  - **Impossible** to run on large datasets as ImageNet

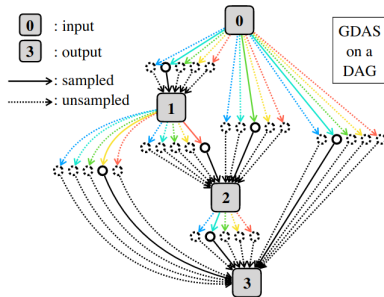
A lot of research aims to fix this issue:

# Issues – Memory constraints

- DARTS keeps the **entire one-shot model in memory**, together with its computed tensors
  - This constrains the search space size and the fidelity used to train the one-shot model
  - **Impossible** to run on large datasets as ImageNet

A lot of research aims to fix this issue:

- **GDAS** [Dong et al, 2019] – samples from a Gumbel Softmax distribution to keep only a single architecture in memory




# Issues – Memory constraints

- DARTS keeps the **entire one-shot model in memory**, together with its computed tensors
  - This constrains the search space size and the fidelity used to train the one-shot model
  - **Impossible** to run on large datasets as ImageNet

A lot of research aims to fix this issue:

- **GDAS** [Dong et al, 2019] – samples from a Gumbel Softmax distribution to keep only a single architecture in memory
- **ProxylessNAS** [Cai et al, 2019] – computes approximate gradients on  $\alpha$  keeping only 2 edges between two nodes in memory at a time




# Issues – Memory constraints

- DARTS keeps the **entire one-shot model in memory**, together with its computed tensors
  - This constrains the search space size and the fidelity used to train the one-shot model
  - **Impossible** to run on large datasets as ImageNet

A lot of research aims to fix this issue:

- **GDAS** [Dong et al, 2019] – samples from a Gumbel Softmax distribution to keep only a single architecture in memory
- **ProxylessNAS** [Cai et al, 2019] – computes approximate gradients on  $\alpha$  keeping only 2 edges between two nodes in memory at a time
- **PC-DARTS** [Xu et al, 2020] – performs the search on a subset of the channels in the one-shot model



## Questions to Answer for Yourself / Discuss with Friends

- Repetition:  
What is the main difference between DARTS and the other one-shot NAS methods we saw before?
- Repetition:  
How does DARTS optimize the architectural weights and one-shot weights?
- Repetition:  
What are DARTS' main issues and how can they be fixed?
- Discussion:  
RobustDARTS stops the architecture optimization early, before the curvature of the validation loss becomes high. Why do you think this works?  
[Hint: think about the discretization step after the DARTS search.]