

# AutoML: Dynamic Configuration & Learning

## Dynamic Configuration

Bernd Bischl   Frank Hutter   Lars Kotthoff  
Marius Lindauer   Joaquin Vanschoren

# Iterative Optimization Heuristics

- Many iterative heuristics in algorithms are dynamic and adaptive
  - ① the algorithm's behavior changes over time
  - ② the algorithm's behavior changes based on internal statistics

# Iterative Optimization Heuristics

- Many iterative heuristics in algorithms are dynamic and adaptive
  - ① the algorithm's behavior changes over time
  - ② the algorithm's behavior changes based on internal statistics
- These heuristics might control other hyperparameters of the algorithms

# Iterative Optimization Heuristics

- Many iterative heuristics in algorithms are dynamic and adaptive
  - ① the algorithm's behavior changes over time
  - ② the algorithm's behavior changes based on internal statistics
- These heuristics might control other hyperparameters of the algorithms
- Example: learning rate schedules for training DNNs
  - ① exponential decaying learning rate: based on number of iterations, learning rate decreases

# Iterative Optimization Heuristics

- Many iterative heuristics in algorithms are dynamic and adaptive
  - ① the algorithm's behavior changes over time
  - ② the algorithm's behavior changes based on internal statistics
- These heuristics might control other hyperparameters of the algorithms
- Example: learning rate schedules for training DNNs
  - ① exponential decaying learning rate: based on number of iterations, learning rate decreases
  - ② Reduce learning rate on plateaus: if the learning stagnates for some time, the learning rate is decreased by a factor

# Iterative Optimization Heuristics

- Many iterative heuristics in algorithms are dynamic and adaptive
  - ① the algorithm's behavior changes over time
  - ② the algorithm's behavior changes based on internal statistics
- These heuristics might control other hyperparameters of the algorithms
- Example: learning rate schedules for training DNNs
  - ① exponential decaying learning rate: based on number of iterations, learning rate decreases
  - ② Reduce learning rate on plateaus: if the learning stagnates for some time, the learning rate is decreased by a factor
- other examples: restart probability of search, mutation rate of evolutionary algorithms, ...

# Parametrization of Learning Rate Schedules

- How can we parameterize learning rate schedules?
  - ① exponential decaying learning rate:
    - ★ initial learning rate
    - ★ minimal learning rate
    - ★ multiplicative factor

# Parametrization of Learning Rate Schedules

- How can we parameterize learning rate schedules?

- ① exponential decaying learning rate:

- ★ initial learning rate
    - ★ minimal learning rate
    - ★ multiplicative factor

- ② Reduce learning rate on plateaus:

- ★ patience (in number of epochs)
    - ★ patience threshold
    - ★ decreasing factor
    - ★ cool-down break (in number of epochs)



# Parametrization of Learning Rate Schedules

- How can we parameterize learning rate schedules?

- ① exponential decaying learning rate:

- ★ initial learning rate
    - ★ minimal learning rate
    - ★ multiplicative factor

- ② Reduce learning rate on plateaus:

- ★ patience (in number of epochs)
    - ★ patience threshold
    - ★ decreasing factor
    - ★ cool-down break (in number of epochs)

~> Many hyperparameters only to control a single hyperparameter

# Parametrization of Learning Rate Schedules

- How can we parameterize learning rate schedules?

- ① exponential decaying learning rate:

- ★ initial learning rate
    - ★ minimal learning rate
    - ★ multiplicative factor

- ② Reduce learning rate on plateaus:

- ★ patience (in number of epochs)
    - ★ patience threshold
    - ★ decreasing factor
    - ★ cool-down break (in number of epochs)

~> Many hyperparameters only to control a single hyperparameter

- Still not guaranteed that optimal setting of e.g. learning rate schedules will lead to optimal learning behavior
  - ▶ Learning rate schedules are only heuristics

# Dynamic Algorithm Configuration

- So far, we assumed that an algorithm runs with static settings
- However, settings, such as learning rate, have to be adapted over time

## Definition

Let

- $\lambda \in \Lambda$  be a hyperparameter configuration of an algorithm  $\mathcal{A}$ ,

# Dynamic Algorithm Configuration

- So far, we assumed that an algorithm runs with static settings
- However, settings, such as learning rate, have to be adapted over time

## Definition

Let

- $\lambda \in \Lambda$  be a hyperparameter configuration of an algorithm  $\mathcal{A}$ ,
- $p(\mathcal{D})$  be a probability distribution over meta datasets  $\mathcal{D} \in \mathbf{D}$ ,

# Dynamic Algorithm Configuration

- So far, we assumed that an algorithm runs with static settings
- However, settings, such as learning rate, have to be adapted over time

## Definition

Let

- $\lambda \in \Lambda$  be a hyperparameter configuration of an algorithm  $\mathcal{A}$ ,
- $p(\mathcal{D})$  be a probability distribution over meta datasets  $\mathcal{D} \in \mathbf{D}$ ,
- $s^{(t)}$  be a state description of  $\mathcal{A}$  solving  $\mathcal{D}$  at time point  $t$ ,

# Dynamic Algorithm Configuration

- So far, we assumed that an algorithm runs with static settings
- However, settings, such as learning rate, have to be adapted over time

## Definition

Let

- $\lambda \in \Lambda$  be a hyperparameter configuration of an algorithm  $\mathcal{A}$ ,
- $p(\mathcal{D})$  be a probability distribution over meta datasets  $\mathcal{D} \in \mathbf{D}$ ,
- $s^{(t)}$  be a state description of  $\mathcal{A}$  solving  $\mathcal{D}$  at time point  $t$ ,
- $c : \Pi \times \mathbf{D} \rightarrow \mathbb{R}$  be a cost metric assessing the cost of a control policy  $\pi \in \Pi$  on  $\mathcal{D} \in \mathbf{D}$

# Dynamic Algorithm Configuration

- So far, we assumed that an algorithm runs with static settings
- However, settings, such as learning rate, have to be adapted over time

## Definition

Let

- $\lambda \in \Lambda$  be a hyperparameter configuration of an algorithm  $\mathcal{A}$ ,
- $p(\mathcal{D})$  be a probability distribution over meta datasets  $\mathcal{D} \in \mathbf{D}$ ,
- $s^{(t)}$  be a state description of  $\mathcal{A}$  solving  $\mathcal{D}$  at time point  $t$ ,
- $c : \Pi \times \mathbf{D} \rightarrow \mathbb{R}$  be a cost metric assessing the cost of a control policy  $\pi \in \Pi$  on  $\mathcal{D} \in \mathbf{D}$

the *dynamic algorithm configuration problem* is to obtain a policy  $\pi^* : s_t \times \mathcal{D} \mapsto \lambda$  by optimizing its cost across a distribution of datasets:

$$\pi^* \in \arg \min_{\pi \in \Pi} \int_{\mathbf{D}} p(\mathcal{D}) c(\pi, \mathcal{D}) \, d\mathcal{D}$$

# Dynamic Algorithm Configuration as Contextual MDP [Biedenkapp et al. 2020]

State  $s^{(t)}$  are described by statistics gathered in the algorithm run



# Dynamic Algorithm Configuration as Contextual MDP [Biedenkapp et al. 2020]

State  $s^{(t)}$  are described by statistics gathered in the algorithm run

Action  $a^{(t)}$  change hyperparameters according to some control policy  $\pi$

# Dynamic Algorithm Configuration as Contextual MDP [Biedenkapp et al. 2020]

- State  $s^{(t)}$  are described by statistics gathered in the algorithm run
- Action  $a^{(t)}$  change hyperparameters according to some control policy  $\pi$
- Transition run the algorithm from state  $s^{(t)}$  to  $s^{(t+1)}$  for a "short" moment by using the hyperparameters defined by  $a^{(t)}$

# Dynamic Algorithm Configuration as Contextual MDP [Biedenkapp et al. 2020]

State  $s^{(t)}$  are described by statistics gathered in the algorithm run

Action  $a^{(t)}$  change hyperparameters according to some control policy  $\pi$

Transition run the algorithm from state  $s^{(t)}$  to  $s^{(t+1)}$  for a "short" moment by using the hyperparameters defined by  $a^{(t)}$

Reward  $r^{(t)}$  Return your current solution quality (or an approximation)

# Dynamic Algorithm Configuration as Contextual MDP [Biedenkapp et al. 2020]

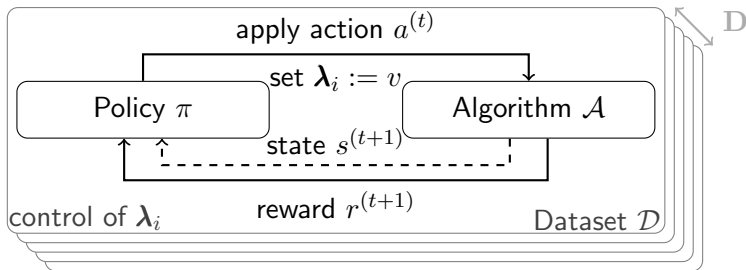
State  $s^{(t)}$  are described by statistics gathered in the algorithm run

Action  $a^{(t)}$  change hyperparameters according to some control policy  $\pi$

Transition run the algorithm from state  $s^{(t)}$  to  $s^{(t+1)}$  for a "short" moment by using the hyperparameters defined by  $a^{(t)}$

Reward  $r^{(t)}$  Return your current solution quality (or an approximation)

Context  $\mathcal{D}$  A given dataset (or task)



# Solving Dynamic Algorithm Configuration

Solve unknown MDP by using reinforcement learning (RL):

$$\mathcal{V}_{\mathcal{D}}^{\pi}(s^{(t)}) = \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k r_{\mathcal{D}}^{(t+k+1)} \mid s^{(t)} = s \right]$$

# Solving Dynamic Algorithm Configuration

Solve unknown MDP by using reinforcement learning (RL):

$$\begin{aligned}\mathcal{V}_{\mathcal{D}}^{\pi}(s^{(t)}) &= \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k r_{\mathcal{D}}^{(t+k+1)} | s^{(t)} = s \right] \\ &= \mathbb{E} \left[ r_{\mathcal{D}}^{(t+1)} + \gamma \mathcal{V}_{\mathcal{D}}^{\pi}(s^{(t+1)}) | s^{(t+1)} \sim \mathcal{T}_{\mathcal{D}}(s^{(t)}, \pi(s^{(t)})) \right]\end{aligned}$$

# Solving Dynamic Algorithm Configuration

Solve unknown MDP by using reinforcement learning (RL):

$$\mathcal{V}_{\mathcal{D}}^{\pi}(s^{(t)}) = \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k r_{\mathcal{D}}^{(t+k+1)} | s^{(t)} = s \right]$$

$$= \mathbb{E} \left[ r_{\mathcal{D}}^{(t+1)} + \gamma \mathcal{V}_{\mathcal{D}}^{\pi}(s^{(t+1)}) | s^{(t+1)} \sim \mathcal{T}_{\mathcal{D}}(s^{(t)}, \pi(s^{(t)})) \right]$$

$$\pi^* \in \arg \max_{\pi \in \Pi} \int_{\mathbf{D}} p(\mathcal{D}) \int_{\mathcal{S}^{(0)}} \Pr(s^{(0)}) \cdot \mathcal{V}_{\mathcal{D}}^{\pi}(s^{(0)}) \, ds^{(0)} \, d\mathcal{D}$$

# Solving Dynamic Algorithm Configuration

Solve unknown MDP by using reinforcement learning (RL):

$$\mathcal{V}_{\mathcal{D}}^{\pi}(s^{(t)}) = \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k r_{\mathcal{D}}^{(t+k+1)} \mid s^{(t)} = s \right]$$

$$= \mathbb{E} \left[ r_{\mathcal{D}}^{(t+1)} + \gamma \mathcal{V}_{\mathcal{D}}^{\pi}(s^{(t+1)}) \mid s^{(t+1)} \sim \mathcal{T}_{\mathcal{D}}(s^{(t)}, \pi(s^{(t)})) \right]$$

$$\pi^* \in \arg \max_{\pi \in \Pi} \int_{\mathbf{D}} p(\mathcal{D}) \int_{\mathcal{S}^{(0)}} \Pr(s^{(0)}) \cdot \mathcal{V}_{\mathcal{D}}^{\pi}(s^{(0)}) \, ds^{(0)} \, d\mathcal{D}$$

$\rightsquigarrow$  equivalent to Dynamic Algorithm Configuration definition



- Challenge: Evaluating a policies on all datasets is often not feasible

# Dynamic Algorithm Configuration across Datasets [Biedenkapp et al. 2020]

- Challenge: Evaluating a policies on all datasets is often not feasible
- Curriculum learning [Bengio et al. 2009] showed that we should have a curriculum of tasks we tackle

- Challenge: Evaluating a policies on all datasets is often not feasible
- Curriculum learning [Bengio et al. 2009] showed that we should have a curriculum of tasks we tackle
- Self-paced learning [Kumar et al. 2010] tries to automatically find such as a curriculum
  - ▶ Focus on "easy" tasks where the agent can improve most:

- Challenge: Evaluating a policies on all datasets is often not feasible
- Curriculum learning [Bengio et al. 2009] showed that we should have a curriculum of tasks we tackle
- Self-paced learning [Kumar et al. 2010] tries to automatically find such as a curriculum
  - ▶ Focus on "easy" tasks where the agent can improve most:

$$\max_{\pi, \mathbf{v}} \mathcal{C}(\pi, \mathbf{v}, K) = \sum_{i=1}^{|\mathbf{D}|} \mathbf{v}_i \mathcal{R}_i(\pi) - \frac{1}{K} \sum_{i=1}^{|\mathbf{D}|} \mathbf{v}_i$$

with  $\theta$  being the agent's policy parameters and  $\mathbf{v}$  being a masking vector for choosing the tasks at hand.