

AutoML: Beyond AutoML

Racing for Algorithm Configuration

Bernd Bischl Frank Hutter Lars Kotthoff
Marius Lindauer Joaquin Vanschoren

State-of-the-art Algorithm Configuration

SMAC: Sequential Model-based Algorithm Configuration [Hutter et al. 2011]

- Bayesian Optimization +
- aggressive racing +
- adaptive capping (for optimizing runtime)

Algorithm 1 SMAC

Input : instance set \mathcal{I} , Algorithm \mathcal{A} with configuration space Λ , Initial configuration λ_0 , performance metric c , Configuration budget b

run history $\mathcal{D}_{\text{Hist}} \leftarrow$ initial design based on λ_0 ;

// $\mathcal{D}_{\text{Hist}} = (\lambda, i, c(i, \lambda))_i$

while b remains **do**

|

Algorithm 2 SMAC

Input : instance set \mathcal{I} , Algorithm \mathcal{A} with configuration space Λ , Initial configuration λ_0 , performance metric c , Configuration budget b

run history $\mathcal{D}_{\text{Hist}} \leftarrow$ initial design based on λ_0 ; $// \mathcal{D}_{\text{Hist}} = (\lambda, i, c(i, \lambda))_i$

while b remains **do**

$\hat{c} \leftarrow$ train empirical performance model based on run history $\mathcal{D}_{\text{Hist}}$;

Algorithm 3 SMAC

Input : instance set \mathcal{I} , Algorithm \mathcal{A} with configuration space Λ , Initial configuration λ_0 , performance metric c , Configuration budget b

run history $\mathcal{D}_{\text{Hist}} \leftarrow$ initial design based on λ_0 ; // $\mathcal{D}_{\text{Hist}} = (\lambda, i, c(i, \lambda))_i$

while b remains **do**

$\hat{c} \leftarrow$ train empirical performance model based on run history $\mathcal{D}_{\text{Hist}}$;

$\Lambda_{\text{challengers}} \leftarrow$ select configurations based on \hat{c} ;

Algorithm 4 SMAC

Input : instance set \mathcal{I} , Algorithm \mathcal{A} with configuration space Λ , Initial configuration λ_0 , performance metric c , Configuration budget b

run history $\mathcal{D}_{\text{Hist}} \leftarrow$ initial design based on λ_0 ; // $\mathcal{D}_{\text{Hist}} = (\lambda, i, c(i, \lambda))_i$

while b remains **do**

$\hat{c} \leftarrow$ train empirical performance model based on run history $\mathcal{D}_{\text{Hist}}$;

$\Lambda_{\text{challengers}} \leftarrow$ select configurations based on \hat{c} ;

$\hat{\lambda}, \mathcal{D}_{\text{Hist}} \leftarrow$ intensify($\Lambda_{\text{challengers}}, \hat{\lambda}$); // racing and capping

Algorithm 5 SMAC

Input : instance set \mathcal{I} , Algorithm \mathcal{A} with configuration space Λ , Initial configuration λ_0 , performance metric c , Configuration budget b

run history $\mathcal{D}_{\text{Hist}} \leftarrow$ initial design based on λ_0 ; // $\mathcal{D}_{\text{Hist}} = (\lambda, i, c(i, \lambda))_i$

while b remains **do**

$\hat{c} \leftarrow$ train empirical performance model based on run history $\mathcal{D}_{\text{Hist}}$;

$\Lambda_{\text{challengers}} \leftarrow$ select configurations based on \hat{c} ;

$\hat{\lambda}, \mathcal{D}_{\text{Hist}} \leftarrow$ intensify($\Lambda_{\text{challengers}}, \hat{\lambda}$); // racing and capping

return $\hat{\lambda}$

Comparisons on N instances [Hutter et al. 2009]

- **Basic(N)** uses a pretty basic comparison: $better_N(\lambda', \lambda)$:
 - ▶ Compare λ' and λ based on N instances

Comparisons on N instances [Hutter et al. 2009]

- **Basic(N)** uses a pretty basic comparison: $better_N(\lambda', \lambda)$:
 - ▶ Compare λ' and λ based on N instances
 - ▶ How does this relate to cross-validation?

- **Basic(N)** uses a pretty basic comparison: $better_N(\lambda', \lambda)$:
 - ▶ Compare λ' and λ based on N instances
 - ▶ How does this relate to cross-validation?
- Problem: How to set N ? Problems of large N ? Small N ?

- **Basic(N)** uses a pretty basic comparison: $better_N(\lambda', \lambda)$:
 - ▶ Compare λ' and λ based on N instances
 - ▶ How does this relate to cross-validation?
- Problem: How to set N ? Problems of large N ? Small N ?
 - ▶ Problem of large N : evaluations are slow
 - ▶ Problem of small N : overfitting to a small set of instances
 - ↪ Tradeoff: Choose N of moderate size

Comparisons on N instances [Hutter et al. 2009]

Question: Which N instances should we use?

- 1 N different instances for each configuration
- 2 The same set of N instances for the entire run

Comparisons on N instances [Hutter et al. 2009]

Question: Which N instances should we use?

- 1 N different instances for each configuration
- 2 The same set of N instances for the entire run

Answer: the same N instances, so that we compare apples with apples
(but: using the same instances can also yield overtuning)

Comparisons on N instances [Hutter et al. 2009]

Question: Which N instances should we use?

- 1 N different instances for each configuration
- 2 The same set of N instances for the entire run

Answer: the same N instances, so that we compare apples with apples
(but: using the same instances can also yield overtuning)

If we sampled different instances for each configuration:

- Some configurations would randomly get easier instances
- Those configurations would look better than they really are

Question: For randomized algorithms, how should we set the seeds?

- 1 Sample a new seed for each algorithm run
- 2 Fix the seeds together with the instances

Comparisons on N instances [Hutter et al. 2009]

Question: For randomized algorithms, how should we set the seeds?

- 1 Sample a new seed for each algorithm run
- 2 Fix the seeds together with the instances

Answer: just like for instances, fix them to compare apples to apples

Comparisons on N instances [Hutter et al. 2009]

Question: For randomized algorithms, how should we set the seeds?

- 1 Sample a new seed for each algorithm run
- 2 Fix the seeds together with the instances

Answer: just like for instances, fix them to compare apples to apples

In summary, for each run of Basic(N):

pick N (instance, seed) pairs and use them for evaluating each λ .

Comparisons on N instances [Hutter et al. 2009]

Question: For randomized algorithms, how should we set the seeds?

- 1 Sample a new seed for each algorithm run
- 2 Fix the seeds together with the instances

Answer: just like for instances, fix them to compare apples to apples

In summary, for each run of Basic(N):

pick N (instance, seed) pairs and use them for evaluating each λ .
(Different Basic(N) runs can use different instances and seeds.)

The concept of overtuning

Very related to overfitting in machine learning

- Performance improves on the training set
- Performance does not improve on the test set, and may even degrade

The concept of overtuning

Very related to overfitting in machine learning

- Performance improves on the training set
- Performance does not improve on the test set, and may even degrade

More pronounced for heterogeneous benchmark sets


- But it even happens for very homogeneous sets
- Indeed, one can even overfit on a single instance, to the [seeds](#) used for training

Overtuning Visualized

- Example: minimizing SLS solver runlengths for a single SAT instance
- Training cost, e.g., with $N=100$:
average runlengths across 100 runs with different seeds
- Test cost of $\hat{\lambda}$ here based on 1000 new seeds

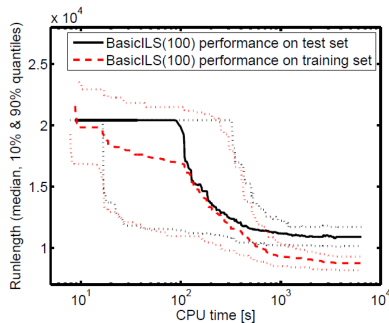
Overtuning Visualized

- Example: minimizing SLS solver runlengths for a single SAT instance
- Training cost, e.g., with $N=100$:
average runlengths across 100 runs with different seeds
- Test cost of $\hat{\lambda}$ here based on 1000 new seeds



Overtuning Visualized

- Example: minimizing SLS solver runlengths for a single SAT instance
- Training cost, e.g., with $N=100$:
average runlengths across 100 runs with different seeds
- Test cost of $\hat{\lambda}$ here based on 1000 new seeds




Basic(N) Test Results with Various N

- Example: minimizing SLS solver runlengths for a single SAT instance
- Training cost, e.g., with $N=?$:
average runlengths across N runs with different seeds
- Test cost of $\hat{\lambda}$ here based on 1000 new seeds


Basic(N) Test Results with Various N

- Example: minimizing SLS solver runlengths for a single SAT instance
- **Training cost**, e.g., with $N=?$:
average runlengths across N runs with different seeds
- **Test cost** of $\hat{\lambda}$ here based on 1000 new seeds



Basic(N) Test Results with Various N


- Example: minimizing SLS solver runlengths for a single SAT instance
- **Training cost**, e.g., with $N=?$:
average runlengths across N runs with different seeds
- **Test cost** of $\hat{\lambda}$ here based on 1000 new seeds



Which of these results corresponds to $N = 1$,
 $N = 10$, and $N = 100$?

Basic(N) Test Results with Various N

- Example: minimizing SLS solver runlengths for a single SAT instance
- Training cost, e.g., with $N=?$:
average runlengths across N runs with different seeds
- Test cost of $\hat{\lambda}$ here based on 1000 new seeds




Which of these results corresponds to $N = 1$, $N = 10$, and $N = 100$?

- 1 $N=1$: blue, $N=10$: red, $N=100$ dashed black
- 2 $N=1$: dashed black, $N=10$: red, $N=100$ blue

Basic(N) Test Results with Various N

- Example: minimizing SLS solver runlengths for a single SAT instance
- Training cost, e.g., with $N=?$:
average runlengths across N runs with different seeds
- Test cost of $\hat{\lambda}$ here based on 1000 new seeds



Which of these results corresponds to $N = 1$, $N = 10$, and $N = 100$?

- 1 $N=1$: blue, $N=10$: red, $N=100$ dashed black
- 2 $N=1$: dashed black, $N=10$: red, $N=100$ blue

Correct Answer: 1

Aggressive Racing (inspired by FocusedILS) [Hutter et al. 2009]

Intuition: get the best of both worlds

- Perform more runs for good configurations
 - to avoid overtuning
- Quickly reject poor configurations
 - to make progress more quickly

Aggressive Racing (inspired by FocusedILS) [Hutter et al. 2009]

Intuition: get the best of both worlds

- Perform more runs for good configurations
 - to avoid overtuning
- Quickly reject poor configurations
 - to make progress more quickly

Definition: $N(\lambda)$ and $c_N(\lambda)$

$N(\lambda)$ denotes the number of runs executed for λ so far.

$\hat{c}_N(\lambda)$ denotes the cost estimate of λ based on N runs.

Aggressive Racing (inspired by FocusedILS) [Hutter et al. 2009]

Intuition: get the best of both worlds

- Perform more runs for good configurations
 - to avoid overtuning
- Quickly reject poor configurations
 - to make progress more quickly

Definition: $N(\lambda)$ and $c_N(\lambda)$

$N(\lambda)$ denotes the number of runs executed for λ so far.

$\hat{c}_N(\lambda)$ denotes the cost estimate of λ based on N runs.

In the beginning: $N(\lambda) = 0$ for every configuration λ

Definition: domination

$\lambda^{(1)}$ dominates $\lambda^{(2)}$ if

- $N(\lambda^{(1)}) \geq N(\lambda^{(2)})$ and
- $\hat{c}_{N(\lambda^{(2)})}(\lambda^{(1)}) \leq \hat{c}_{N(\lambda^{(2)})}(\lambda^{(2)})$.

I.e.: we have at least as many runs for $\lambda^{(1)}$ and its cost is at least as low.

Aggressive Racing (inspired by FocusedILS) [Hutter et al. 2009]

Definition: domination

$\lambda^{(1)}$ dominates $\lambda^{(2)}$ if

- $N(\lambda^{(1)}) \geq N(\lambda^{(2)})$ and
- $\hat{c}_{N(\lambda^{(2)})}(\lambda^{(1)}) \leq \hat{c}_{N(\lambda^{(2)})}(\lambda^{(2)})$.

I.e.: we have at least as many runs for $\lambda^{(1)}$ and its cost is at least as low.

better($\lambda', \hat{\lambda}$) in a nutshell

- $\hat{\lambda}$ is the current configuration to beat (incumbent)

Aggressive Racing (inspired by FocusedILS) [Hutter et al. 2009]

Definition: domination

$\lambda^{(1)}$ dominates $\lambda^{(2)}$ if

- $N(\lambda^{(1)}) \geq N(\lambda^{(2)})$ and
- $\hat{c}_{N(\lambda^{(2)})}(\lambda^{(1)}) \leq \hat{c}_{N(\lambda^{(2)})}(\lambda^{(2)})$.

I.e.: we have at least as many runs for $\lambda^{(1)}$ and its cost is at least as low.

better($\lambda', \hat{\lambda}$) in a nutshell

- $\hat{\lambda}$ is the current configuration to beat (incumbent)
- Perform runs of λ' until either
 - ▶ $\hat{\lambda}$ dominates $\lambda' \rightsquigarrow$ reject λ' , or
 - ▶ λ' dominates $\hat{\lambda} \rightsquigarrow$ change current incumbent ($\hat{\lambda} \leftarrow \lambda'$)

Aggressive Racing (inspired by FocusedILS) [Hutter et al. 2009]

Definition: domination

$\lambda^{(1)}$ dominates $\lambda^{(2)}$ if

- $N(\lambda^{(1)}) \geq N(\lambda^{(2)})$ and
- $\hat{c}_{N(\lambda^{(2)})}(\lambda^{(1)}) \leq \hat{c}_{N(\lambda^{(2)})}(\lambda^{(2)})$.

I.e.: we have at least as many runs for $\lambda^{(1)}$ and its cost is at least as low.

better($\lambda', \hat{\lambda}$) in a nutshell

- $\hat{\lambda}$ is the current configuration to beat (incumbent)
- Perform runs of λ' until either
 - ▶ $\hat{\lambda}$ dominates $\lambda' \rightsquigarrow$ reject λ' , or
 - ▶ λ' dominates $\hat{\lambda} \rightsquigarrow$ change current incumbent ($\hat{\lambda} \leftarrow \lambda'$)
- Over time: perform extra runs of $\hat{\lambda}$ to gain more confidence in it

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'			
λ''			

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'	2		

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'	2	10	

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'	2	10	
\rightarrow reject, since $\hat{c}_2(\lambda') = 6 > \hat{c}_2(\hat{\lambda}) = 2.5$			

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'	2	10	
\rightarrow reject, since $\hat{c}_2(\lambda') = 6 > \hat{c}_2(\hat{\lambda}) = 2.5$			
λ''	3		

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'	2	10	
\rightarrow reject, since $\hat{c}_2(\lambda') = 6 > \hat{c}_2(\hat{\lambda}) = 2.5$			
λ''	3	1	

Toy Example

- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'	2	10	
\rightarrow reject, since $\hat{c}_2(\lambda') = 6 > \hat{c}_2(\hat{\lambda}) = 2.5$			
λ''	3	1	5

Toy Example


- Let $\hat{\lambda}$ be the incumbent (evaluated on $i^{(1)}, i^{(2)}, i^{(3)}$)
- We'll look at challengers λ' and λ''

	$i^{(1)}$	$i^{(2)}$	$i^{(3)}$
$\hat{\lambda}$	3	2	10
λ'	2	10	
\rightarrow reject, since $\hat{c}_2(\lambda') = 6 > \hat{c}_2(\hat{\lambda}) = 2.5$			
λ''	3	1	5

- new incumbent: $\hat{\lambda} \leftarrow \lambda''$
- Perform an additional run for new $\hat{\lambda}$ to increase confidence over time


Racing achieves the best of both worlds

Aggressive racing (aka FocusedILS): Fast progress and no overtuning



Racing achieves the best of both worlds

Aggressive racing (aka FocusedILS): Fast progress and no overtuning



Overview of Racing

Input : candidate configurations Λ_{new} , cutoff κ_{max} , previously evaluated runs \mathcal{D}_{Hist} , budget T , incumbent $\hat{\lambda}$
while Λ_{new} *not empty* **do**
 $\lambda^{(t)} \leftarrow \text{getNext}(\Lambda_{new});$

Overview of Racing

Input : candidate configurations Λ_{new} , cutoff κ_{max} , previously evaluated runs \mathcal{D}_{Hist} , budget T , incumbent $\hat{\lambda}$

while Λ_{new} *not empty* **do**

- $\lambda^{(t)} \leftarrow \text{getNext}(\Lambda_{new});$
- $i, s \leftarrow$ instance and seed drawn uniformly at random;
- $c \leftarrow \text{EvaluateRun}(\hat{\lambda}, i, s, \kappa_{max});$
- $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\hat{\lambda}, i, s, c);$

Overview of Racing

Input : candidate configurations Λ_{new} , cutoff κ_{max} , previously evaluated runs \mathcal{D}_{Hist} , budget T , incumbent $\hat{\lambda}$

while Λ_{new} *not empty* **do**

- $\lambda^{(t)} \leftarrow \text{getNext}(\Lambda_{new});$
- $i, s \leftarrow$ instance and seed drawn uniformly at random;
- $c \leftarrow \text{EvaluateRun}(\hat{\lambda}, i, s, \kappa_{max});$
- $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\hat{\lambda}, i, s, c);$
- while** *true* **do**
 - $\mathcal{I}^+, s^+ \leftarrow \text{getAlreadyEvaluatedOn}(\hat{\lambda}, \mathcal{D}_{Hist});$
 - $\mathcal{I}^{(t)}, s^{(t)} \leftarrow \text{getAlreadyEvaluatedOn}(\lambda^{(t)}, \mathcal{D}_{Hist});$
 - $i^{(t)}, s^{(t)} \leftarrow$ drawn uniformly at random from $\mathcal{I}^+ \setminus \mathcal{I}^{(t)}$ and $s^+ \setminus s^{(t)};$

Overview of Racing

Input : candidate configurations Λ_{new} , cutoff κ_{max} , previously evaluated runs \mathcal{D}_{Hist} , budget T , incumbent $\hat{\lambda}$

while Λ_{new} not empty **do**

$\lambda^{(t)} \leftarrow \text{getNext}(\Lambda_{new});$

$i, s \leftarrow$ instance and seed drawn uniformly at random;

$c \leftarrow \text{EvaluateRun}(\hat{\lambda}, i, s, \kappa_{max});$

$\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\hat{\lambda}, i, s, c);$

while true **do**

$\mathcal{I}^+, s^+ \leftarrow \text{getAlreadyEvaluatedOn}(\hat{\lambda}, \mathcal{D}_{Hist});$

$\mathcal{I}^{(t)}, s^{(t)} \leftarrow \text{getAlreadyEvaluatedOn}(\lambda^{(t)}, \mathcal{D}_{Hist});$

$i^{(t)}, s^{(t)} \leftarrow$ drawn uniformly at random from $\mathcal{I}^+ \setminus \mathcal{I}^{(t)}$ and $s^+ \setminus s^{(t)};$

$c_i \leftarrow \text{EvaluateRun}(\lambda^{(t)}, i^{(t)}, s^{(t)}, \kappa_{max});$

$\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\lambda^{(t)}, i^{(t)}, s^{(t)}, c^{(t)});$

Overview of Racing

Input : candidate configurations Λ_{new} , cutoff κ_{max} , previously evaluated runs \mathcal{D}_{Hist} , budget T , incumbent $\hat{\lambda}$

while Λ_{new} *not empty* **do**

- $\lambda^{(t)} \leftarrow \text{getNext}(\Lambda_{new});$
- $i, s \leftarrow$ instance and seed drawn uniformly at random;
- $c \leftarrow \text{EvaluateRun}(\hat{\lambda}, i, s, \kappa_{max});$
- $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\hat{\lambda}, i, s, c);$
- while** *true* **do**
 - $\mathcal{I}^+, s^+ \leftarrow \text{getAlreadyEvaluatedOn}(\hat{\lambda}, \mathcal{D}_{Hist});$
 - $\mathcal{I}^{(t)}, s^{(t)} \leftarrow \text{getAlreadyEvaluatedOn}(\lambda^{(t)}, \mathcal{D}_{Hist});$
 - $i^{(t)}, s^{(t)} \leftarrow$ drawn uniformly at random from $\mathcal{I}^+ \setminus \mathcal{I}^{(t)}$ and $s^+ \setminus s^{(t)};$
 - $c_i \leftarrow \text{EvaluateRun}(\lambda^{(t)}, i^{(t)}, s^{(t)}, \kappa_{max});$
 - $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\lambda^{(t)}, i^{(t)}, s^{(t)}, c^{(t)});$
 - if** *average cost of $\lambda^{(t)}$ > average cost of $\hat{\lambda}$ across $\mathcal{I}^{(t)}$ and $s^{(t)}$* **then**
 - \perp break;

Overview of Racing

Input : candidate configurations Λ_{new} , cutoff κ_{max} , previously evaluated runs \mathcal{D}_{Hist} , budget T , incumbent $\hat{\lambda}$

while Λ_{new} *not empty* **do**

- $\lambda^{(t)} \leftarrow \text{getNext}(\Lambda_{new});$
- $i, s \leftarrow$ instance and seed drawn uniformly at random;
- $c \leftarrow \text{EvaluateRun}(\hat{\lambda}, i, s, \kappa_{max});$
- $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\hat{\lambda}, i, s, c);$
- while** *true* **do**
 - $\mathcal{I}^+, s^+ \leftarrow \text{getAlreadyEvaluatedOn}(\hat{\lambda}, \mathcal{D}_{Hist});$
 - $\mathcal{I}^{(t)}, s^{(t)} \leftarrow \text{getAlreadyEvaluatedOn}(\lambda^{(t)}, \mathcal{D}_{Hist});$
 - $i^{(t)}, s^{(t)} \leftarrow$ drawn uniformly at random from $\mathcal{I}^+ \setminus \mathcal{I}^{(t)}$ and $s^+ \setminus s^{(t)};$
 - $c_i \leftarrow \text{EvaluateRun}(\lambda^{(t)}, i^{(t)}, s^{(t)}, \kappa_{max});$
 - $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\lambda^{(t)}, i^{(t)}, s^{(t)}, c^{(t)});$
 - if** *average cost of $\lambda^{(t)}$ > average cost of $\hat{\lambda}$ across $\mathcal{I}^{(t)}$ and $s^{(t)}$* **then**
 - \perp break;
 - else if** *average cost of $\lambda^{(t)}$ < average cost of $\hat{\lambda}$ and $\mathcal{I}^+ = \mathcal{I}^{(t)}$ and $s^+ = s^{(t)}$* **then**
 - $\perp \hat{\lambda} \leftarrow \lambda^{(t)};$

Overview of Racing

Input : candidate configurations Λ_{new} , cutoff κ_{max} , previously evaluated runs \mathcal{D}_{Hist} , budget T , incumbent $\hat{\lambda}$

while Λ_{new} *not empty* **do**

- $\lambda^{(t)} \leftarrow \text{getNext}(\Lambda_{new});$
- $i, s \leftarrow$ instance and seed drawn uniformly at random;
- $c \leftarrow \text{EvaluateRun}(\hat{\lambda}, i, s, \kappa_{max});$
- $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\hat{\lambda}, i, s, c);$
- while** *true* **do**
 - $\mathcal{I}^+, s^+ \leftarrow \text{getAlreadyEvaluatedOn}(\hat{\lambda}, \mathcal{D}_{Hist});$
 - $\mathcal{I}^{(t)}, s^{(t)} \leftarrow \text{getAlreadyEvaluatedOn}(\lambda^{(t)}, \mathcal{D}_{Hist});$
 - $i^{(t)}, s^{(t)} \leftarrow$ drawn uniformly at random from $\mathcal{I}^+ \setminus \mathcal{I}^{(t)}$ and $s^+ \setminus s^{(t)};$
 - $c_i \leftarrow \text{EvaluateRun}(\lambda^{(t)}, i^{(t)}, s^{(t)}, \kappa_{max});$
 - $\mathcal{D}_{Hist} \leftarrow \mathcal{D}_{Hist} \cup (\lambda^{(t)}, i^{(t)}, s^{(t)}, c^{(t)});$
 - if** *average cost of $\lambda^{(t)}$ > average cost of $\hat{\lambda}$ across $\mathcal{I}^{(t)}$ and $s^{(t)}$* **then**
 - \perp break;
 - else if** *average cost of $\lambda^{(t)}$ < average cost of $\hat{\lambda}$ and $\mathcal{I}^+ = \mathcal{I}^{(t)}$ and $s^+ = s^{(t)}$* **then**
 - $\perp \hat{\lambda} \leftarrow \lambda^{(t)};$

if *time spent exceeds T or Λ_{new} is empty* **then**

- \perp **return** $\hat{\lambda}, \mathcal{D}_{Hist}$